

# Querying Series

In this lecture, we'll talk about one of the primary data types of the Pandas library, the Series. You'll learn about the structure of the Series, how to query and merge Series objects together, and the importance of thinking about parallelization when engaging in data science programming.

## loc & iloc( ) Attributes ¶

In [3]:

```
# A pandas Series can be queried either by the index position or the index label. If you don't give an
# index to the series when querying, the position and the label are effectively the same values. To
# query by numeric location, starting at zero, use the iloc attribute. To query by the index label,
# you can use the loc attribute.

# Lets start with an example. We'll use students enrolled in classes coming from a dictionary
import pandas as pd
students_classes = {'Alice': 'Physics',
                    'Jack': 'Chemistry',
                    'Molly': 'English',
                    'Sam': 'History'}
s = pd.Series(students_classes)
s
```

Out[3]:

```
Alice      Physics
Jack       Chemistry
Molly      English
Sam        History
dtype: object
```

In [3]:

```
# So, for this series, if you wanted to see the fourth entry we would use the iloc
# attribute with the parameter 3.
s.iloc[3]
```

Out[3]:

```
'History'
```

In [3]:

```
# If you wanted to see what class Molly has, we would use the loc attribute with  
a parameter  
# of Molly.  
s.loc['Molly']
```

Out[3]:

'English'

Keep in mind that `iloc` and `loc` are not methods, they **are attributes**. So you don't use parentheses to query them, but **square brackets** instead, which is called the indexing operator. In Python this calls get or set for an item depending on the context of its use.

This might seem a bit confusing if you're used to languages where encapsulation of attributes, variables, and properties is common, such as in Java.

## Shorthands & When You Shouldn't Use Them

Pandas tries to make our code a bit more readable and provides a sort of smart syntax using the indexing operator directly on the series itself.

In [5]:

```
# Pandas tries to make our code a bit more readable and provides a sort of smart  
syntax using  
# the indexing operator directly on the series itself. For instance, if you pass  
in an integer parameter,  
# the operator will behave as if you want it to query via the iloc attribute  
s[3]
```

Out[5]:

'History'

In [6]:

```
# If you pass in an object, it will query as if you wanted to use the label base  
d loc attribute.  
s['Molly']
```

Out[6]:

'English'

So what happens if your index is a list of integers? This is a bit complicated and Pandas can't determine automatically whether you're intending to query by index position or index label. So you need to be careful when using the indexing operator on the Series itself.

The safer option is to be more explicit and **use the `iloc` or `loc` attributes directly**.

In [5]:

```
# Here's an example using class and their classcode information, where classes are indexed by
# classcodes, in the form of integers
class_code = {99: 'Physics',
              100: 'Chemistry',
              101: 'English',
              102: 'History'}
s = pd.Series(class_code)
```

In [7]:

```
# If we try and call s[0] we get a key error because there's no item in the classes list with
# an index of zero, instead we have to call iloc explicitly if we want the first item.
try:
    print ("Trying s[0]")
    s[0]
except:
    print ("KeyError")
    print(s.iloc[0])
```

```
Trying s[0]
KeyError
Physics
```

So, that didn't call `s.iloc[0]` underneath as one might expect, instead it generates an error

Now we know how to get data out of the series, let's talk about working with the data.

## Data Manipulation

A common task is to want to consider all of the values inside of a series and do some sort of operation. This could be trying to find a certain number, or summarizing data or transforming the data in some way.

In [8]:

```
# A typical programmatic approach to this would be to iterate over all the items in the series,
# and invoke the operation one is interested in. For instance, we could create a Series of
# integers representing student grades, and just try and get an average grade

grades = pd.Series([90, 80, 70, 60])

total = 0
for grade in grades:
    total+=grade
print(total/len(grades))
```

75.0

# Vectorisation in Pandas

This works, but it's slow. Modern computers can do many tasks simultaneously, especially, but not only, tasks involving mathematics.

Pandas and the underlying numpy libraries support a method of computation called **vectorization**. Vectorization works with most of the functions in the numpy library, including the sum function.

In [9]:

```
# Here's how we would really write the code using the numpy sum method. First we need to import the numpy module

import numpy as np

# Then we just call np.sum and pass in an iterable item. In this case, our panda series.

total = np.sum(grades)
print(total/len(grades))
```

75.0

## Magic Functions: Example of Comparing RunTimes

Note: The `head()` function look at the top five items in a series.

In [11]:

```
# Now both of these methods create the same value, but is one actually faster? The Jupyter Notebook has a magic function which can help.

# First, let's create a big series of random numbers. This is used a lot when demonstrating techniques with Pandas
numbers = pd.Series(np.random.randint(0,1000,10000))

# Now lets look at the top five items in that series to make sure they actually seem random.
numbers.head()
```

Out[11]:

```
0    288
1    796
2    170
3    156
4    253
dtype: int64
```

In [12]:

```
# We can actually verify that length of the series is correct using the len function
len(numbers)
```

Out[12]:

10000

Ok, we're confident now that we have a big series. The ipython interpreter has something called magic functions begin with a percentage sign. If we type this sign and then hit the Tab key, you can see a list of the available magic functions. You could write your own magic functions too, but that's a little bit outside of the scope of this course.

## %%timeit Function

Here, we're actually going to use what's called a cellular magic function. These start with two percentage signs and wrap the code in the current Jupyter cell. The function we're going to use is called timeit. This function will run our code a few times to determine, on average, how long it takes.

Let's run timeit with our original iterative code. You can give timeit the number of loops that you would like to run. By default, it is 1,000 loops. I'll ask timeit here to use 100 runs because we're recording this. Note that in order to use a cellular magic function, it has to be **the first line of the cell**

In [15]:

```
%%timeit -n 100
total = 0
for number in numbers:
    total+=number

total/len(numbers)
```

1.09 ms  $\pm$  10  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 100 loops each)

Not bad. Timeit ran the code and it doesn't seem to take very long at all. Now let's try with vectorization.

In [12]:

```
%%timeit -n 100
total = np.sum(numbers)
total/len(numbers)
```

66.4  $\mu$ s  $\pm$  3  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 100 loops each)

Wow! This is a pretty shocking difference in the speed and demonstrates why one should be aware of parallel computing features and start thinking in functional programming terms. Put more simply, **vectorization is the ability for a computer to execute multiple instructions at once**, and with high performance chips, especially graphics cards, you can get dramatic speedups. Modern graphics cards can run thousands of instructions in parallel.

# Broadcasting

A Related feature in pandas and nummy is called broadcasting. **With broadcasting, you can apply an operation to every value in the series, changing the series.** For instance, if we wanted to increase every random variable by 2, we could do so quickly using the += operator directly on the Series object.

In [13]:

```
# Let's look at the head of our series  
numbers.head()
```

Out[13]:

```
0    288  
1    796  
2    170  
3    156  
4    253  
dtype: int64
```

In [14]:

```
# And now lets just increase everything in the series by 2  
numbers+=2  
numbers.head()
```

Out[14]:

```
0    290  
1    798  
2    172  
3    158  
4    255  
dtype: int64
```

The procedural way of doing this would be to iterate through all of the items in the series and increase the values directly. Pandas does support iterating through a series much like a dictionary, allowing you to unpack values easily.

In [15]:

```
# We can use the iteritems() function which returns a label and value  
for label, value in numbers.iteritems():  
    # now for the item which is returned, lets call set_value()  
    numbers.set_value(label, value+2)  
# And we can check the result of this computation  
numbers.head()
```

Out[15]:

```
0    292  
1    800  
2    174  
3    160  
4    257  
dtype: int64
```

So the result is the same, though you may notice a warning depending upon the version of pandas being used. But if you find yourself iterating pretty much *any time* in pandas, you should question whether you're doing things in the best possible way.

Lets take a look at some speed comparisons. First, lets try five loops using the iterative approach

In [16]:

```
%%timeit -n 10
# we'll create a blank new series of items to deal with
s = pd.Series(np.random.randint(0,1000,1000))
# And we'll just rewrite our loop from above.
for label, value in s.iteritems():
    s.loc[label]= value+2
```

124 ms ± 1.46 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

Now lets try that using the broadcasting methods

In [24]:

```
%%timeit -n 10
# We need to recreate a series
s = pd.Series(np.random.randint(0,1000,1000))
# And we just broadcast with +=
s+=2
```

The slowest run took 33.58 times longer than the fastest. This could mean that an intermediate result is being cached.  
1.29 ms ± 2.55 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

Amazing. Not only is it significantly faster, but it's more concise and even easier to read too. The typical mathematical operations you would expect are vectorized, and the **numpy** documentation outlines what it takes to create vectorized functions of your own.

## Multiple Data Types in Series

One last note on using the indexing operators to access series data. The `.loc` attribute lets you not only modify data in place, but also add new data as well. If the value you pass in as the index doesn't exist, then a new entry is added. And keep in mind, indices can have mixed types. While it's important to be aware of the typing going on underneath, Pandas will automatically change the underlying NumPy types as appropriate.

In [26]:

```
# Here's an example using a Series of a few numbers.
s = pd.Series([1, 2, 3])

# We could add some new value, maybe a university course
s.loc['History'] = 102

s
```

Out[26]:

```
0          1
1          2
2          3
History    102
dtype: int64
```

## Non-Unique Index Values

We see that mixed types for data values or index labels are no problem for Pandas. Since "History" is not in the original list of indices, `s.loc['History']` essentially creates a new element in the series, with the index named "History", and the value of 102.

Up until now I've shown only examples of a series where the index values were unique. I want to end this lecture by showing an example where index values are not unique, and this makes pandas Series a **little different** conceptually then, for instance, a **relational database**.

In [19]:

```
# Lets create a Series with students and the courses which they have taken
students_classes = pd.Series({'Alice': 'Physics',
                              'Jack': 'Chemistry',
                              'Molly': 'English',
                              'Sam': 'History'})

students_classes
```

Out[19]:

```
Alice      Physics
Jack       Chemistry
Molly      English
Sam        History
dtype: object
```



In [20]:

```
# Now lets create a Series just for some new student Kelly, which lists all of the courses
# she has taken. We'll set the index to Kelly, and the data to be the names of courses.
kelly_classes = pd.Series(['Philosophy', 'Arts', 'Math'], index=['Kelly', 'Kelly', 'Kelly'])
kelly_classes
```

Out[20]:

```
Kelly    Philosophy
Kelly           Arts
Kelly           Math
dtype: object
```

## Appending a Series to Another Series using .append()

In [21]:

```
# Finally, we can append all of the data in this new Series to the first using the .append()
# function.
all_students_classes = students_classes.append(kelly_classes)

# This creates a series which has our original people in it as well as all of Kelly's courses
all_students_classes
```

Out[21]:

```
Alice      Physics
Jack       Chemistry
Molly      English
Sam        History
Kelly      Philosophy
Kelly           Arts
Kelly           Math
dtype: object
```

There are a couple of important considerations when using append. First, Pandas will take the series and try to infer the best data types to use. In this example, everything is a string, so there's no problems here. Second, the append method doesn't actually change the underlying Series objects, it instead returns a new series which is made up of the two appended together. This is a common pattern in pandas - by default returning a new object instead of modifying in place - and one you should come to expect. By printing the original series we can see that that series hasn't changed.

In [22]:

```
students_classes
```

Out[22]:

```
Alice      Physics
Jack       Chemistry
Molly      English
Sam        History
dtype: object
```

In [23]:

```
# Finally, we see that when we query the appended series for Kelly, we don't get
a single value,
# but a series itself.
all_students_classes.loc['Kelly'] ## WE GET A SERIES!!
```

Out[23]:

```
Kelly      Philosophy
Kelly      Arts
Kelly      Math
dtype: object
```

## Summary

In this lecture, we focused on one of the primary data types of the Pandas library, the Series. You learned how to query the Series, with `.loc` and `.iloc`, that the Series is an indexed data structure, how to merge two Series objects together with `append()`, and the importance of vectorization.

## Pandas DataFrame

There are many more methods associated with the Series object that we haven't talked about. But with these basics down, we'll move on to talking about the **Panda's two-dimensional data structure**, the DataFrame. The DataFrame is very similar to the series object, but includes multiple columns of data, and is the structure that you'll spend the majority of your time working with when cleaning and aggregating data.