

03-05 Pivot Table

A pivot table is a way of summarizing data in a DataFrame for a particular purpose. **It makes heavy use of the aggregation function.**

A pivot table is itself a DataFrame, where the rows represent **one variable** that you're interested in, **the columns another**, and the cell's some aggregate value. A pivot table also tends to **includes marginal values** as well, which are the sums for each column and row. This allows you to be able to see the relationship between two variables at just a glance.

Creating Pivot Tables in Pandas

In [1]:

```
# Lets take a look at pivot tables in pandas
import pandas as pd
import numpy as np
```

In [2]:

```
# Here we have the Times Higher Education World University Ranking dataset, which is one of the most influential university measures. Let's import the dataset and see what it looks like
df = pd.read_csv('datasets/cwurData.csv')
df.head()
```

Out[2]:

	world_rank	institution	country	national_rank	quality_of_education	alumni_employer
0	1	Harvard University	USA	1	7	
1	2	Massachusetts Institute of Technology	USA	2	9	1
2	3	Stanford University	USA	3	17	1
3	4	University of Cambridge	United Kingdom	1	10	2
4	5	California Institute of Technology	USA	4	2	2

Here we can see each institution's rank, country, quality of education, other metrics, and overall score. Let's say we want to create a new column called Rank_Level, where institutions with world ranking 1-100 are categorized as first tier and those with world ranking 101 - 200 are second tier, ranking 201 - 300 are third tier, after 301 is other top universities.

In [3]:

```
def create_category(ranking):
    # Since the rank is just an integer, I'll just do a bunch of if/elif statements
    if (ranking >= 1) & (ranking <= 100):
        return "First Tier Top University"
    elif (ranking >= 101) & (ranking <= 200):
        return "Second Tier Top University"
    elif (ranking >= 201) & (ranking <= 300):
        return "Third Tier Top University"
    return "Other Top University"

# Now we can apply this to a single column of data to create a new series
df['Rank_Level'] = df['world_rank'].apply(lambda x: create_category(x))
# And lets look at the result
df.head()
```

Out[3]:

	world_rank	institution	country	national_rank	quality_of_education	alumni_employer
0	1	Harvard University	USA	1	7	
1	2	Massachusetts Institute of Technology	USA	2	9	1
2	3	Stanford University	USA	3	17	1
3	4	University of Cambridge	United Kingdom	1	10	2
4	5	California Institute of Technology	USA	4	2	2

In [4]:

```
# A pivot table allows us to pivot out one of these columns a new column headers
and compare it against
# another column as row indices. Let's say we want to compare rank level versus
country of the universities
# and we want to compare in terms of overall score

# To do this, we tell Pandas we want the values to be Score, and index to be the
country and the columns to be
# the rank levels. Then we specify that the aggregation function, and here we'll
use the NumPy mean to get the
# average rating for universities in that country

df.pivot_table(values='score', index='country', columns='Rank_Level', aggfunc=[n
p.mean]).head()
```

Out[4]:

Rank_Level	mean			
	First Tier Top University	Other Top University	Second Tier Top University	Third Tier Top University
country				
Argentina	NaN	44.672857	NaN	NaN
Australia	47.9425	44.645750	49.2425	47.285000
Austria	NaN	44.864286	NaN	47.066667
Belgium	51.8750	45.081000	49.0840	46.746667
Brazil	NaN	44.499706	49.5650	NaN

We can see a hierarchical dataframe where the index, or rows, are by country and the columns have two levels, the top level indicating that the mean value is being used and the second level being our ranks. In this example we only have one variable, the mean, that we are looking at, so we don't really need a heirarchical index.

We notice that there are some NaN values, for example, the first row, Argentina. The NaN values indicate that Argentina has only observations in the "Other Top Universities" category

Multiple Hierarchical Indices

Now, pivot tables aren't limited to one function that you might want to apply. You can pass a named parameter, `aggfunc`, which is a **list** of the different functions to apply, and pandas will provide you with the result **using hierarchical column names**. Let's try that same query, but pass in the `max()` function too.

This example should make *Hierarchical Column Index* clearer.

In [5]:

```
df.pivot_table(values='score', index='country', columns='Rank_Level', aggfunc=[np.mean, np.max]).head()
```

Out[5]:

Rank_Level	mean				amax			
	First Tier Top University	Other Top University	Second Tier Top University	Third Tier Top University	First Tier Top University	Other Top University	Second Tier Top University	Third Tier Top University
country								
Argentina	NaN	44.672857	NaN	NaN	NaN	45.66	NaN	1
Australia	47.9425	44.645750	49.2425	47.285000	51.61	45.97	50.40	47
Austria	NaN	44.864286	NaN	47.066667	NaN	46.29	NaN	47
Belgium	51.8750	45.081000	49.0840	46.746667	52.03	46.21	49.73	47
Brazil	NaN	44.499706	49.5650	NaN	NaN	46.08	49.82	1

Marginal Values

In [6]:

```
df.pivot_table(values='score', index='country', columns='Rank_Level', aggfunc=[np.mean, np.max], margins=True).head()
```

Out[6]:

Rank_Level	mean					All	amax		
	First Tier Top University	Other Top University	Second Tier Top University	Third Tier Top University	First Tier Top University		Other Top University	Second Tier Top University	
country									
Argentina	NaN	44.672857	NaN	NaN	44.672857	NaN	45.66		
Australia	47.9425	44.645750	49.2425	47.285000	45.825517	51.61	45.97	5	
Austria	NaN	44.864286	NaN	47.066667	45.139583	NaN	46.29		
Belgium	51.8750	45.081000	49.0840	46.746667	47.011000	52.03	46.21	4	
Brazil	NaN	44.499706	49.5650	NaN	44.781111	NaN	46.08	4	

Pivot Tables as MultiLevel DataFrames

A pivot table is just a multi-level dataframe, and we can access series or cells in the dataframe in a similar way as we do so for a regular dataframe. In the below example, we can see the columns are hierarchical. The top level column indices have two categories: mean and max, and the lower level column indices have four categories, which are the four rank levels.

In [7]:

```
# Let's create a new dataframe from our previous example
new_df=df.pivot_table(values='score', index='country', columns='Rank_Level', agg
func=[np.mean, np.max],
                        margins=True)
# Now let's look at the index
print(new_df.index)
# And let's look at the columns
print(new_df.columns)
```

```
Index(['Argentina', 'Australia', 'Austria', 'Belgium', 'Brazil', 'Bulgaria',
      'Canada', 'Chile', 'China', 'Colombia', 'Croatia', 'Cyprus',
      'Czech Republic', 'Denmark', 'Egypt', 'Estonia', 'Finland',
      'France',
      'Germany', 'Greece', 'Hong Kong', 'Hungary', 'Iceland', 'India',
      'Iran',
      'Ireland', 'Israel', 'Italy', 'Japan', 'Lebanon', 'Lithuania',
      'Malaysia', 'Mexico', 'Netherlands', 'New Zealand', 'Norway',
      'Poland',
      'Portugal', 'Puerto Rico', 'Romania', 'Russia', 'Saudi Arabia',
      'Serbia', 'Singapore', 'Slovak Republic', 'Slovenia', 'South Africa',
      'South Korea', 'Spain', 'Sweden', 'Switzerland', 'Taiwan', 'Thailand',
      'Turkey', 'USA', 'Uganda', 'United Arab Emirates', 'United Kingdom',
      'Uruguay', 'All'],
      dtype='object', name='country')
MultiIndex([('mean', 'First Tier Top University'),
            ('mean', 'Other Top University'),
            ('mean', 'Second Tier Top University'),
            ('mean', 'Third Tier Top University'),
            ('mean', 'All'),
            ('amax', 'First Tier Top University'),
            ('amax', 'Other Top University'),
            ('amax', 'Second Tier Top University'),
            ('amax', 'Third Tier Top University'),
            ('amax', 'All')],
          names=[None, 'Rank_Level'])
```

How would we query this if we want to get the average scores of First Tier Top University levels in each country? We would just need to **make two dataframe projections**, the first for the mean, then the second for the top tier

In [8]:

```
new_df['mean']['First Tier Top University'].head()
```

Out[8]:

```
country
Argentina      NaN
Australia      47.9425
Austria         NaN
Belgium        51.8750
Brazil          NaN
Name: First Tier Top University, dtype: float64
```

In [9]:

```
# We can see that the output is a series object which we can confirm by printing
the type. Remember that when
# you project a single column of values out of a DataFrame you get a series.
type(new_df['mean']['First Tier Top University'])
```

Out[9]:

```
pandas.core.series.Series
```

In [10]:

```
# What if we want to find the country that has the maximum average score on First
Tier Top University level?
# We can use the idxmax() function.
print(type(new_df['mean']['First Tier Top University'].idxmax()))
new_df['mean']['First Tier Top University'].idxmax()
```

```
<class 'str'>
```

Out[10]:

```
'United Kingdom'
```

Now, the `idxmax()` function isn't special for pivot tables, it's a built in function to the Series object. We don't have time to go over all pandas functions and attributes, and I want to encourage you to explore the API to learn more deeply what is available to you.

Stacking & Unstacking

To achieve a different shape of your pivot table, you can do so with the `stack()` and `unstack()` functions.

- **Stacking** is pivoting the **lowermost** column **index** to become the **innermost row index**.
- **Unstacking** is the opposite of stacking, pivoting the **innermost** row index to become the **lowermost** column index.

In [11]:

```
# If you want to achieve a different shape of your pivot table, you can do so with the stack and unstack functions. Stacking is pivoting the lowermost column index to become the innermost row index. Unstacking is the inverse of stacking, pivoting the innermost row index to become the lowermost column index. An example will help make this clear

# Let's look at our pivot table first to refresh what it looks like
new_df.head()
```

Out[11]:

Rank_Level	mean					amax		
	First Tier Top University	Other Top University	Second Tier Top University	Third Tier Top University	All	First Tier Top University	Other Top University	Second Tier Top University
country								
Argentina	NaN	44.672857	NaN	NaN	44.672857	NaN	45.66	
Australia	47.9425	44.645750	49.2425	47.285000	45.825517	51.61	45.97	5
Austria	NaN	44.864286	NaN	47.066667	45.139583	NaN	46.29	
Belgium	51.8750	45.081000	49.0840	46.746667	47.011000	52.03	46.21	4
Brazil	NaN	44.499706	49.5650	NaN	44.781111	NaN	46.08	4

In [12]:

```
# Now let's try stacking, this should move the lowermost column, so the tiers of the university rankings, to the inner most row
new_df=new_df.stack()
new_df.head()
```

Out[12]:

country	Rank_Level	mean	amax
Argentina	Other Top University	44.672857	45.66
	All	44.672857	45.66
Australia	First Tier Top University	47.942500	51.61
	Other Top University	44.645750	45.97
	Second Tier Top University	49.242500	50.40

In [13]:

```
# In the original pivot table, rank levels are the lowermost column, after stacking, rank levels become the innermost index, appearing to the right after country

# Now let's try unstacking
new_df.unstack().head()
```

Out[13]:

Rank_Level	mean					amax		
	First Tier Top University	Other Top University	Second Tier Top University	Third Tier Top University	All	First Tier Top University	Other Top University	Second Tier Top University
country								
Argentina	NaN	44.672857	NaN	NaN	44.672857	NaN	45.66	
Australia	47.9425	44.645750	49.2425	47.285000	45.825517	51.61	45.97	5
Austria	NaN	44.864286	NaN	47.066667	45.139583	NaN	46.29	
Belgium	51.8750	45.081000	49.0840	46.746667	47.011000	52.03	46.21	4
Brazil	NaN	44.499706	49.5650	NaN	44.781111	NaN	46.08	4

In [14]:

```
# That seems to restore our dataframe to its original shape. What do you think would happen if we unstacked twice in a row?
new_df.unstack().unstack().head()
```

Out[14]:

```
Rank_Level
mean  First Tier Top University  country
Argentina  NaN
Australia  47.9425
Austria    NaN
Belgium    51.8750
Brazil     NaN

dtype: float64
```

In [15]:

```
# We actually end up unstacking all the way to just a single column, so a series object is returned. This column is just a "value", the meaning of which is denoted by the hierarchical index of operation, rank, and country.
```

Summary

So that's pivot tables. This has been a pretty short description, but they're incredibly useful when dealing with numeric data, especially if you're trying to summarize the data in some form. You'll regularly be creating new pivot tables on slices of data, whether you're exploring the data yourself or preparing data for others to report on. And of course, you can pass any function you want to the aggregate function, including those that you define yourself.

In []: