

## 03-06 Date / Time Functionality

In today's lecture, where we'll be looking at the time series and date functionality in pandas. Manipulating dates and time is quite flexible in Pandas and thus allows us to conduct more analysis such as time series analysis, which we will talk about soon. Actually, pandas was originally created by Wes McKinney to handle date and time data when he worked as a consultant for hedge funds.

In [1]:

```
# Let's bring in pandas and numpy as usual
import pandas as pd
import numpy as np
```

### Time Classes

- **Timestamp** - This represents a single timestamp and associates values with points in time.
- **DatetimeIndex** - This is the index of timestamp.
- **Period** - Helps us classify spans of time - like a month, or a year.
- **PeriodIndex** - This is the index of a period.

### Time Classes: Timestamp

Timestamp is interchangeable with Python's datetime in most cases. It is also flexible and is able to **modify different input types** and output the datetime in a standardised format.

Timestamp also has useful **attributes** such as `isoweekday()` which shows what day of the week it was or it will be, and can extract elements of a specific timestamp as shown below.

In [2]:

```
# Pandas has four main time related classes. Timestamp, DatetimeIndex, Period, and PeriodIndex. First, let's
# look at Timestamp. It represents a single timestamp and associates values with
# points in time.

# For example, let's create a timestamp using a string 9/1/2019 10:05AM, and here we have our timestamp.
# Timestamp is interchangeable with Python's datetime in most cases.
pd.Timestamp('9/1/2019 10:05AM')
```

Out[2]:

```
Timestamp('2019-09-01 10:05:00')
```

In [3]:

```
# We can also create a timestamp by passing multiple parameters such as year, month, date, hour, minute, separately
pd.Timestamp(2019, 12, 20, 0, 0, 5)
```

Out[3]:

```
Timestamp('2019-12-20 00:00:05')
```

In [4]:

```
# Timestamp also has some useful attributes, such as isoweekday(), which shows the weekday of the timestamp
# note that 1 represents Monday and 7 represents Sunday
pd.Timestamp(2019, 12, 20, 0, 0).isoweekday()
```

Out[4]:

```
5
```

In [5]:

```
# You can find extract the specific year, month, day, hour, minute, second from a timestamp
pd.Timestamp(2019, 12, 20, 5, 2, 23).second
```

Out[5]:

```
23
```

## Time Classes: Period

Often we aren't interested in a specific point in time but instead a span of time. This is where the Period class comes into play. Period represents a single time span, such as a specific day or month.

Periods are sensitive to its **granularity** (ie. the finest grained piece we provided). So if we define a period by its month and year, the period is by *month*. If we define a period by its date, month and year, the period is by *day*.

**Arithmetic** can also be applied to periods and the period object **encapsulates the granularity for arithmetic**.

In [6]:

```
# Suppose we weren't interested in a specific point in time and instead wanted a span of time. This is where the Period class comes into play. Period represents a single time span, such as a specific day or month.

# Here we are creating a period that is January 2016,
pd.Period('1/2016')
```

Out[6]:

```
Period('2016-01', 'M')
```

In [7]:

```
# You'll notice when we print that out that the granularity of the period is M f
or month, since that was the
# finest grained piece we provided. Here's an example of a period that is March
5th, 2016.
pd.Period('3/5/2016')
```

Out[7]:

```
Period('2016-03-05', 'D')
```

In [8]:

```
# Period objects represent the full timespan that you specify. Arithmetic on per
iod is very easy and
# intuitive, for instance, if we want to find out 5 months after January 2016, w
e simply plus 5
pd.Period('1/2016') + 5
```

Out[8]:

```
Period('2016-06', 'M')
```

In [9]:

```
# From the result, you can see we get June 2016. If we want to find out two days
before March 5th 2016, we
# simply subtract 2
pd.Period('1/1/2016') - 2
```

Out[9]:

```
Period('2015-12-30', 'D')
```

## TimeStamp: DatetimeIndex and PeriodIndex

In [10]:

```
# The index of a timestamp is DatetimeIndex. Let's look at a quick example. Firs
t, let's create our example
# series t1, we'll use the Timestamp of September 1st, 2nd and 3rd of 2016. When
we look at the series, each
# Timestamp is the index and has a value associated with it, in this case, a, b
and c.
```

```
t1 = pd.Series(list('abc'), [pd.Timestamp('2016-09-01'), pd.Timestamp('2016-09-0
2'),
                        pd.Timestamp('2016-09-03')])
t1
```

Out[10]:

```
2016-09-01    a
2016-09-02    b
2016-09-03    c
dtype: object
```

In [11]:

```
# Looking at the type of our series index, we see that it's DatetimeIndex.  
type(t1.index)
```

Out[11]:

```
pandas.core.indexes.datetimes.DatetimeIndex
```

In [12]:

```
# Similarly, we can create a period-based index as well.  
t2 = pd.Series(list('def'), [pd.Period('2016-09'), pd.Period('2016-10'),  
                             pd.Period('2016-11')])  
t2
```

Out[12]:

```
2016-09    d  
2016-10    e  
2016-11    f  
Freq: M, dtype: object
```

In [13]:

```
# Looking at the type of the ts2.index, we can see that it's PeriodIndex.  
type(t2.index)
```

Out[13]:

```
pandas.core.indexes.period.PeriodIndex
```

## Converting to Datetime using `to_datetime()`

In [14]:

```
# Now, let's look into how to convert to Datetime. Suppose we have a list of dates as strings and we want to  
# create a new dataframe  
  
# I'm going to try a bunch of different date formats  
d1 = ['2 June 2013', 'Aug 29, 2014', '2015-06-26', '7/12/16']  
  
# And just some random data  
ts3 = pd.DataFrame(np.random.randint(10, 100, (4,2)), index=d1,  
                   columns=list('ab'))  
ts3
```

Out[14]:

	a	b
2 June 2013	84	46
Aug 29, 2014	75	10
2015-06-26	22	85
7/12/16	60	34

In [15]:

```
# Using pandas to_datetime, pandas will try to convert these to Datetime and put them in a standard format.
```

```
ts3.index = pd.to_datetime(ts3.index)
ts3
```

Out[15]:

	a	b
2013-06-02	84	46
2014-08-29	75	10
2015-06-26	22	85
2016-07-12	60	34

In [16]:

```
# to_datetime() also has options to change the date parse order. For example, we  
# can pass in the argument dayfirst = True to parse the date in European date.
```

```
pd.to_datetime('4.7.12', dayfirst=True)
```

Out[16]:

```
Timestamp('2012-07-04 00:00:00')
```

## Timedelta

Timedeltas are differences in times. This is not the same as a a period, but conceptually similar. For instance, if we want to take the **difference between September 3rd and September 1st**, we get a Timedelta of two days.

**Arithmetic operators** on different timestamps are also valid.

In [17]:

```
pd.Timestamp('9/3/2016') - pd.Timestamp('9/1/2016')
```

Out[17]:

```
Timedelta('2 days 00:00:00')
```

In [18]:

```
# We can also do something like find what the date and time is for 12 days and t  
hree hours past September 2nd,  
# at 8:10 AM.
```

```
pd.Timestamp('9/2/2016 8:10AM') + pd.Timedelta('12D 3H')
```

Out[18]:

```
Timestamp('2016-09-14 11:10:00')
```

## Offset

Offset is similar to timedelta, but it follows specific calendar duration rules. Offset allows **flexibility in terms of types of time intervals**. Besides hour, day, week, month, etc it also has business day `.BusinessDay()`, end of month `.MonthEnd()` etc..

In [19]:

```
# Let's create a timestamp, and see what day is that  
# Note that 0 represents Monday, 6 represents Sunday.  
# 4th Sep 2016 is a Sunday.  
pd.Timestamp('9/4/2016').weekday()
```

Out[19]:

6

In [20]:

```
# Now we can now add the timestamp with a week ahead  
# This is code to obtain date of the following Sunday.  
pd.Timestamp('9/4/2016') + pd.offsets.Week()
```

Out[20]:

Timestamp('2016-09-11 00:00:00')

In [21]:

```
# Now let's try to do the month end, then we would have the last day of September  
  
pd.Timestamp('9/4/2016') + pd.offsets.MonthEnd()
```

Out[21]:

Timestamp('2016-09-30 00:00:00')

## Working with Dates in a Dataframe

Using `.date_range()`

In [22]:

```
# Next, let's look at a few tricks for working with dates in a DataFrame. Suppose we want to look at nine measurements, taken bi-weekly, every Sunday, starting in October 2016. Using date_range, we can create this DatetimeIndex. In data_range, we have to either specify the start or end date. If it is not explicitly specified, by default, the date is considered the start date. Then we have to specify number of periods, and a frequency. Here, we set it to "2W-SUN", which means biweekly on Sunday

dates = pd.date_range('10-01-2016', periods=9, freq='2W-SUN')
dates
```

Out[22]:

```
DatetimeIndex(['2016-10-02', '2016-10-16', '2016-10-30', '2016-11-13',
               '2016-11-27', '2016-12-11', '2016-12-25', '2017-01-08',
               '2017-01-22'],
              dtype='datetime64[ns]', freq='2W-SUN')
```

In [23]:

```
# There are many other frequencies that you can specify. For example, you can do business day
pd.date_range('10-01-2016', periods=9, freq='B')
```

Out[23]:

```
DatetimeIndex(['2016-10-03', '2016-10-04', '2016-10-05', '2016-10-06',
               '2016-10-07', '2016-10-10', '2016-10-11', '2016-10-12',
               '2016-10-13'],
              dtype='datetime64[ns]', freq='B')
```

In [24]:

```
# Or you can do quarterly, with the quarter start in June
pd.date_range('04-01-2016', periods=12, freq='QS-JUN')
```

Out[24]:

```
DatetimeIndex(['2016-06-01', '2016-09-01', '2016-12-01', '2017-03-01',
               '2017-06-01', '2017-09-01', '2017-12-01', '2018-03-01',
               '2018-06-01', '2018-09-01', '2018-12-01', '2019-03-01'],
              dtype='datetime64[ns]', freq='QS-JUN')
```

In [25]:

```
# Now, let's go back to our weekly on Sunday example and create a DataFrame using these dates, and some random data, and see what we can do with it.

dates = pd.date_range('10-01-2016', periods=9, freq='2W-SUN')
df = pd.DataFrame({'Count 1': 100 + np.random.randint(-5, 10, 9).cumsum(),
                   'Count 2': 120 + np.random.randint(-5, 10, 9)}, index=dates)
df
```

Out[25]:

	Count 1	Count 2
2016-10-02	106	117
2016-10-16	101	124
2016-10-30	104	124
2016-11-13	111	116
2016-11-27	113	127
2016-12-11	113	115
2016-12-25	109	127
2017-01-08	118	123
2017-01-22	115	126

In [26]:

```
# First, we can check what day of the week a specific date is. For example, here we can see that all the dates in our index are on a Sunday. Which matches the frequency that we set
df.index.weekday_name
```

Out[26]:

```
Index(['Sunday', 'Sunday', 'Sunday', 'Sunday', 'Sunday', 'Sunday',
       'Sunday',
       'Sunday', 'Sunday'],
      dtype='object')
```

Using `.diff()`



In [27]:

```
# We can also use diff() to find the difference between each date's value.  
df.diff()
```

Out[27]:

	Count 1	Count 2
2016-10-02	NaN	NaN
2016-10-16	-5.0	7.0
2016-10-30	3.0	0.0
2016-11-13	7.0	-8.0
2016-11-27	2.0	11.0
2016-12-11	0.0	-12.0
2016-12-25	-4.0	12.0
2017-01-08	9.0	-4.0
2017-01-22	-3.0	3.0

## Using .resample()

In [28]:

```
# Suppose we want to know what the mean count is for each month in our DataFrame.  
We can do this using  
# resample. Converting from a higher frequency from a lower frequency is called  
downsampling (we'll talk about  
this in a moment)  
df.resample('M').mean()
```

Out[28]:

	Count 1	Count 2
2016-10-31	103.666667	121.666667
2016-11-30	112.000000	121.500000
2016-12-31	111.000000	121.000000
2017-01-31	116.500000	124.500000

## Using indexing & slicing

In [29]:

```
# Now let's talk about datetime indexing and slicing, which is a wonderful feature of the pandas DataFrame.  
# For instance, we can use partial string indexing to find values from a particular year,  
df['2017']
```

Out[29]:

	Count 1	Count 2
2017-01-08	118	123
2017-01-22	115	126

In [30]:

```
# Or we can do it from a particular month  
df['2016-12']
```

Out[30]:

	Count 1	Count 2
2016-12-11	113	115
2016-12-25	109	127

In [31]:

```
# Or we can even slice on a range of dates For example, here we only want the values from December 2016 onwards.  
df['2016-12':]
```

Out[31]:

	Count 1	Count 2
2016-12-11	113	115
2016-12-25	109	127
2017-01-08	118	123
2017-01-22	115	126

In [32]:

```
df[ '2016' ]
```

Out[32]:

	Count 1	Count 2
2016-10-02	106	117
2016-10-16	101	124
2016-10-30	104	124
2016-11-13	111	116
2016-11-27	113	127
2016-12-11	113	115
2016-12-25	109	127

In [ ]: