
You are currently looking at **version 1.2** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](https://www.coursera.org/learn/python-machine-learning/resources/bANLa) (<https://www.coursera.org/learn/python-machine-learning/resources/bANLa>), course resource.

Assignment 3 - Evaluation

In this assignment you will train several models and evaluate how effectively they predict instances of fraud using data based on [this dataset from Kaggle](https://www.kaggle.com/dalpozz/creditcardfraud) (<https://www.kaggle.com/dalpozz/creditcardfraud>).

Each row in `fraud_data.csv` corresponds to a credit card transaction. Features include confidential variables `v1` through `v28` as well as `Amount` which is the amount of the transaction.

The target is stored in the `class` column, where a value of 1 corresponds to an instance of fraud and 0 corresponds to an instance of not fraud.

In [1]:

```
import numpy as np
import pandas as pd
```

Question 1

Import the data from `fraud_data.csv`. What percentage of the observations in the dataset are instances of fraud?

This function should return a float between 0 and 1.

In [2]:

```
def answer_one():

    # Your code here
    fraud_df = pd.read_csv("fraud_data.csv")
    frauds = fraud_df[fraud_df['Class'] == 1]

    return (float)(len(fraud_df[fraud_df['Class'] == 1]) / len(fraud_df)) # Return your answer
answer_one()
```

Out[2]:

0.016410823768035772

In [3]:

```
fraud_df = pd.read_csv("fraud_data.csv")
fraud_df.head()
```

Out[3]:

	V1	V2	V3	V4	V5	V6	V7	V8
0	1.176563	0.323798	0.536927	1.047002	-0.368652	-0.728586	0.084678	-0.06924
1	0.681109	-3.934776	-3.801827	-1.147468	-0.735540	-0.501097	1.038865	-0.62697
2	1.140729	0.453484	0.247010	2.383132	0.343287	0.432804	0.093380	0.173310
3	-1.107073	-3.298902	-0.184092	-1.795744	2.137564	-1.684992	-2.015606	-0.00718
4	-0.314818	0.866839	-0.124577	-0.627638	2.651762	3.428128	0.194637	0.670674

5 rows × 30 columns

In [5]:

```
# Use X_train, X_test, y_train, y_test for all of the following questions
from sklearn.model_selection import train_test_split

df = pd.read_csv('fraud_data.csv')

X = df.iloc[:, :-1]
y = df.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

Question 2

Using `X_train`, `X_test`, `y_train`, and `y_test` (as defined above), train a dummy classifier that classifies everything as the majority class of the training data. What is the accuracy of this classifier? What is the recall?

This function should return a tuple with two floats, i.e. (accuracy score, recall score).

In [6]:

```
def answer_two():
    from sklearn.dummy import DummyClassifier
    from sklearn.metrics import recall_score

    # Your code here
    dummy_majority = DummyClassifier(strategy = 'most_frequent').fit(X_train, y_train)
    y_majority_predicted = dummy_majority.predict(X_test)

    return (dummy_majority.score(X_test, y_test), recall_score(y_test, y_majority_predicted)) # Return your answer
answer_two()
```

Out[6]:

(0.98525073746312686, 0.0)

Question 3

Using `X_train`, `X_test`, `y_train`, `y_test` (as defined above), train a SVC classifier using the default parameters. What is the accuracy, recall, and precision of this classifier?

This function should return a tuple with three floats, i.e. (accuracy score, recall score, precision score).

In [7]:

```
def answer_three():
    from sklearn.metrics import recall_score, precision_score
    from sklearn.svm import SVC

    # Your code here
    svc_clf = SVC().fit(X_train, y_train)
    y_predicted = svc_clf.predict(X_test)

    return (svc_clf.score(X_test, y_test), recall_score(y_test, y_predicted),
            precision_score(y_test, y_predicted)) # Return your answer
answer_three()
```

Out[7]:

(0.99078171091445433, 0.375, 1.0)

Question 4

Using the SVC classifier with parameters `{ 'C': 1e9, 'gamma': 1e-07 }`, what is the confusion matrix when using a threshold of -220 on the decision function. Use `X_test` and `y_test`.

This function should return a confusion matrix, a 2x2 numpy array with 4 integers.

In [10]:

```
def answer_four():
    from sklearn.metrics import confusion_matrix
    from sklearn.svm import SVC

    # Your code here
    # Find the predicted y scores of confidence of the decision function
    y_scores = SVC(gamma = 10**(-7), C = 10**9).fit(X_train, y_train).decision_f
unction(X_test)

    # Set a threshold -220 to obtain y_predicted
    y_predicted = np.array([1 if x > -220 else 0 for x in y_scores])

    # Pass this y_predicted value into the confusion matrix
    return confusion_matrix(y_test, y_predicted)# Return your answer
answer_four()
```

Out[10]:

```
array([[5320,  24],
       [  14,  66]])
```

Question 5

Train a logistic regression classifier with default parameters using `X_train` and `y_train`.

For the logistic regression classifier, create a precision recall curve and a roc curve using `y_test` and the probability estimates for `X_test` (probability it is fraud).

Looking at the precision recall curve, what is the recall when the precision is 0.75?

Looking at the roc curve, what is the true positive rate when the false positive rate is 0.16?

This function should return a tuple with two floats, i.e. (recall, true positive rate).

In [19]:

```
def plot_prc(prec,rec):
    %matplotlib notebook
    import matplotlib.pyplot as plt
    plt.figure(figsize = (7,5))
    plt.xlim([0.0, 1.01])
    plt.ylim([0.0, 1.01])
    plt.plot(prec, rec, label='Precision-Recall Curve')
    plt.vlines(0.75, 0.0, 1.01) # Line where precision = 0.75
    plt.xlabel('Precision', fontsize=16)
    plt.ylabel('Recall', fontsize=16)
    plt.axes().set_aspect('equal')
    plt.show()

def plot_roc(fpr, tpr):
    %matplotlib notebook
    import matplotlib.pyplot as plt
    plt.figure(figsize = (7,5))
    plt.xlim([-0.01, 1.00])
    plt.ylim([-0.01, 1.01])
    plt.plot(fpr, tpr, label = 'ROC Curve')
    plt.axes().set_aspect('equal')
    plt.vlines(0.16, -0.01, 1.01)
    plt.show()

def answer_five():

    # Your code here
    # Import relevant modules
    from sklearn.linear_model import LogisticRegression
    from sklearn.metrics import precision_recall_curve, roc_curve

    # Create the classifier and train the model
    lr = LogisticRegression().fit(X_train, y_train)
    # Find the y_scores on this dataset.
    y_scores = lr.decision_function(X_test)
    # Get the precision and recall
    precision, recall, _ = precision_recall_curve(y_test, y_scores)
    # Get the false positive rate and true positive rate
    fpr, tpr, _ = roc_curve(y_test, y_scores)
    # Plot the graphs
    #plot_prc(precision, recall)
    #plot_roc(fpr, tpr)

    return (0.823576, 0.94625)# Return your answer
answer_five()
```

Out[19]:

```
(0.823576, 0.94625)
```

Question 6

Perform a grid search over the parameters listed below for a Logistic Regression classifier, using recall for scoring and the default 3-fold cross validation.

```
'penalty': ['l1', 'l2']
```

```
'C':[0.01, 0.1, 1, 10, 100]
```

From `.cv_results_`, create an array of the mean test scores of each parameter combination. i.e.

	l1	l2
0.01	?	?
0.1	?	?
1	?	?
10	?	?
100	?	?

This function should return a 5 by 2 numpy array with 10 floats.

Note: do not return a DataFrame, just the values denoted by '?' above in a numpy array. You might need to reshape your raw result to meet the format we are looking for.

In [23]:

```
def answer_six():
    from sklearn.model_selection import GridSearchCV
    from sklearn.linear_model import LogisticRegression

    # Your code here
    clf = LogisticRegression()
    grid_values = {'penalty':['l1', 'l2'], 'C':[0.01, 0.1, 1, 10, 100]}
    # Create a gridsearch constructor with the logistic regression classifier, p
arameters, scoring and the cross
    # validation parameter.
    gridsearch = GridSearchCV(clf, param_grid = grid_values, scoring = 'recall',
cv = 3)
    # Fit the training data to it
    gridsearch.fit(X_train, y_train)
    # See the decision function scores
    mean_test_scores = gridsearch.cv_results_['mean_test_score'].reshape(5,2)
    return mean_test_scores# Return your answer
answer_six()
```


Out[23]:

```
array([[ 0.66666667,  0.76086957],
       [ 0.80072464,  0.80434783],
       [ 0.8115942 ,  0.8115942 ],
       [ 0.80797101,  0.8115942 ],
       [ 0.80797101,  0.80797101]])
```

In [24]:

```
# Use the following function to help visualize results from the grid search
def GridSearch_Heatmap(scores):
    %matplotlib notebook
    import seaborn as sns
    import matplotlib.pyplot as plt
    plt.figure()
    sns.heatmap(scores.reshape(5,2), xticklabels=['l1','l2'], yticklabels=[0.01,
0.1, 1, 10, 100])
    plt.yticks(rotation=0);

#GridSearch_Heatmap(answer_six())
#1 hour 10 min
```



In []: