

Python's Runtime Environment

The Python Imaging Library (PIL)

The Python Imaging Library, which is known as PIL or PILLOW, is the main library we use in python for dealing with image files. This library is not included with python - it's what's known as a third party library, which means you have to download and install it yourself. In the Coursera system, this has all been done for you. Lets do a little exploring of pillow in the jupyter notebooks.

In [1]:

```
# You'll recall that we import a library using the `import` keyword.  
import PIL
```

In [5]:

```
# Documentation is a big help in learning a library. There exist standards that  
  make this process easier.  
# For example, most libraries let you check their version using the version attr  
  ibute.  
PIL.__version__
```

Out[5]:

```
'5.4.1'
```

In [6]:

```
# Let's figure out how to open an image with `Pillow`. Python provides some built-in functions to help us  
# understand the functions and objects which are available in libraries. For instance, the help function,  
# when called on any object, returns the object's built-in documentation. Lets try it with our new library  
# module, PIL.  
help(PIL)
```

Help on package PIL:

NAME

PIL - Pillow (Fork of the Python Imaging Library)

DESCRIPTION

Pillow is the friendly PIL fork by Alex Clark and Contributors.
<https://github.com/python-pillow/Pillow/>

Pillow is forked from PIL 1.1.7.

PIL is the Python Imaging Library by Fredrik Lundh and Contributors.

Copyright (c) 1999 by Secret Labs AB.

Use PIL.__version__ for this Pillow version.

PIL.VERSION is the old PIL version and will be removed in the future.

;-)

PACKAGE CONTENTS

BdfFontFile
BlpImagePlugin
BmpImagePlugin
BufrStubImagePlugin
ContainerIO
CurImagePlugin
DcxImagePlugin
DdsImagePlugin
EpsImagePlugin
ExifTags
FitsStubImagePlugin
FliImagePlugin
FontFile
FpxImagePlugin
FtexImagePlugin
GbrImagePlugin
GdImageFile
GifImagePlugin
GimpGradientFile
GimpPaletteFile
GribStubImagePlugin
Hdf5StubImagePlugin
IcnsImagePlugin
IcoImagePlugin
ImImagePlugin
Image
ImageChops
ImageCms
ImageColor
ImageDraw
ImageDraw2
ImageEnhance
ImageFile
ImageFilter
ImageFont
ImageGrab
ImageMath
ImageMode
ImageMorph

ImageOps
ImagePalette
ImagePath
ImageQt
ImageSequence
ImageShow
ImageStat
ImageTk
ImageTransform
ImageWin
ImtImagePlugin
IptcImagePlugin
Jpeg2KImagePlugin
JpegImagePlugin
JpegPresets
McIdasImagePlugin
MicImagePlugin
MpegImagePlugin
MpoImagePlugin
MspImagePlugin
OleFileIO
PSDraw
PaletteFile
PalmImagePlugin
PcdImagePlugin
PcfFontFile
PcxImagePlugin
PdfImagePlugin
PdfParser
PixarImagePlugin
PngImagePlugin
PpmImagePlugin
PsdImagePlugin
PyAccess
SgiImagePlugin
SpiderImagePlugin
SunImagePlugin
TarIO
TgaImagePlugin
TiffImagePlugin
TiffTags
WalImageFile
WebPImagePlugin
WmfImagePlugin
XVThumbImagePlugin
XbmImagePlugin
XpmImagePlugin
_binary
_imaging
_imagingft
_imagingmath
_imagingmorph
_imagingtk
_tkinter_finder
_util
_version
features

DATA

PILLOW_VERSION = '5.4.1'
VERSION = '1.1.7'

```
VERSION
5.4.1
```

```
FILE
/opt/conda/lib/python3.7/site-packages/PIL/__init__.py
```

```
In [7]:
```

```
# This shows us that there are a host of classes available to us in the module,
as well as version information
# and even the file, called __init__.py, which has the source code for the modul
e itself. We could look up
# the source code for this in the Jupyter console if we wanted to. These documen
tation standards make it easy
# to poke around an unexplored library.
#
# Python also has a function called dir() which will list the contents of an obj
ect. This is especially useful
# with modules where you might want to see what classes you might interact with.
Lets list the details of
# the PIL module
dir(PIL)
```

```
Out[7]:
```

```
['PILLOW_VERSION',
 'VERSION',
 '__builtins__',
 '__cached__',
 '__doc__',
 '__file__',
 '__loader__',
 '__name__',
 '__package__',
 '__path__',
 '__spec__',
 '__version__',
 '_plugins']
```

In [8]:

```
# At the top of the list, there is something called Image. This sounds like it c  
ould be interesting, so lets  
# import it directly, and run the help command on it.  
from PIL import Image  
help(Image)
```

Help on module PIL.Image in PIL:

NAME

PIL.Image

DESCRIPTION

```
# The Python Imaging Library.
# $Id$
#
# the Image class wrapper
#
# partial release history:
# 1995-09-09 fl    Created
# 1996-03-11 fl    PIL release 0.0 (proof of concept)
# 1996-04-30 fl    PIL release 0.1b1
# 1999-07-28 fl    PIL release 1.0 final
# 2000-06-07 fl    PIL release 1.1
# 2000-10-20 fl    PIL release 1.1.1
# 2001-05-07 fl    PIL release 1.1.2
# 2002-03-15 fl    PIL release 1.1.3
# 2003-05-10 fl    PIL release 1.1.4
# 2005-03-28 fl    PIL release 1.1.5
# 2006-12-02 fl    PIL release 1.1.6
# 2009-11-15 fl    PIL release 1.1.7
#
# Copyright (c) 1997-2009 by Secret Labs AB.  All rights reserve
d.
# Copyright (c) 1995-2009 by Fredrik Lundh.
#
# See the README file for information on usage and redistributio
n.
#
```

CLASSES

```
builtins.Exception(builtins.BaseException)
    DecompressionBombError
builtins.RuntimeWarning(builtins.Warning)
    DecompressionBombWarning
builtins.object
    Image
    ImagePointHandler
    ImageTransformHandler
```

```
class DecompressionBombError(builtins.Exception)
|   Common base class for all non-exit exceptions.
|
|   Method resolution order:
|       DecompressionBombError
|       builtins.Exception
|       builtins.BaseException
|       builtins.object
|
|   Data descriptors defined here:
|
|   __weakref__
|       list of weak references to the object (if defined)
|
|   -----
```

```
-----
|   Methods inherited from builtins.Exception:
```

```

|   __init__(self, /, *args, **kwargs)
|       Initialize self. See help(type(self)) for accurate sign
ature.
|
| -----
|
| Static methods inherited from builtins.Exception:
|
|   __new__(*args, **kwargs) from builtins.type
|       Create and return a new object. See help(type) for accu
rate signature.
|
| -----
|
| Methods inherited from builtins.BaseException:
|
|   __delattr__(self, name, /)
|       Implement delattr(self, name).
|
|   __getattr__(self, name, /)
|       Return getattr(self, name).
|
|   __reduce__(...)
|       Helper for pickle.
|
|   __repr__(self, /)
|       Return repr(self).
|
|   __setattr__(self, name, value, /)
|       Implement setattr(self, name, value).
|
|   __setstate__(...)
|
|   __str__(self, /)
|       Return str(self).
|
|   with_traceback(...)
|       Exception.with_traceback(tb) --
|       set self.__traceback__ to tb and return self.
|
| -----
|
| Data descriptors inherited from builtins.BaseException:
|
|   __cause__
|       exception cause
|
|   __context__
|       exception context
|
|   __dict__
|
|   __suppress_context__
|
|   __traceback__
|
|   args

```

```

class DecompressionBombWarning(builtins.RuntimeWarning)
|   Base class for warnings about dubious runtime behavior.
|

```



```
Method resolution order:
    DecompressionBombWarning
    builtins.RuntimeWarning
    builtins.Warning
    builtins.Exception
    builtins.BaseException
    builtins.object
```

```
Data descriptors defined here:
```

```
__weakref__
    list of weak references to the object (if defined)
```

```
-----
Methods inherited from builtins.RuntimeWarning:
```

```
__init__(self, /, *args, **kwargs)
    Initialize self. See help(type(self)) for accurate sign
```

ature.

```
-----
Static methods inherited from builtins.RuntimeWarning:
```

```
__new__(*args, **kwargs) from builtins.type
    Create and return a new object. See help(type) for accu
```

rate signature.

```
-----
Methods inherited from builtins.BaseException:
```

```
__delattr__(self, name, /)
    Implement delattr(self, name).
```

```
__getattr__(self, name, /)
    Return getattr(self, name).
```

```
__reduce__(...)
    Helper for pickle.
```

```
__repr__(self, /)
    Return repr(self).
```

```
__setattr__(self, name, value, /)
    Implement setattr(self, name, value).
```

```
__setstate__(...)
```

```
__str__(self, /)
    Return str(self).
```

```
with_traceback(...)
    Exception.with_traceback(tb) --
    set self.__traceback__ to tb and return self.
```

```
-----
Data descriptors inherited from builtins.BaseException:
```

```

__cause__
    exception cause

__context__
    exception context

__dict__

__suppress_context__

__traceback__

args

```

```

class Image(builtins.object)
|   This class represents an image object.  To create
|   :py:class:`~PIL.Image.Image` objects, use the appropriate fa
actory
|   functions.  There's hardly ever any reason to call the Image
constructor
|   directly.

|   * :py:func:`~PIL.Image.open`
|   * :py:func:`~PIL.Image.new`
|   * :py:func:`~PIL.Image.frombytes`

|   Methods defined here:

|   __copy__ = copy(self)

|   __del__(self)

|   __enter__(self)
|       # Context manager support

|   __eq__(self, other)
|       Return self==value.

|   __exit__(self, *args)

|   __getstate__(self)

|   __init__(self)
|       Initialize self.  See help(type(self)) for accurate sign
ature.

|   __ne__(self, other)
|       Return self!=value.

|   __repr__(self)
|       Return repr(self).

|   __setstate__(self, state)

|   alpha_composite(self, im, dest=(0, 0), source=(0, 0))
|       'In-place' analog of Image.alpha_composite. Composites a
n image
|       onto this image.

|       :param im: image to composite over this one
|       :param dest: Optional 2 tuple (left, top) specifying the

```

upper
left corner in this (destination) image.
:param source: Optional 2 (left, top) tuple for the upper left
corner in the overlay source image, or 4 tuple (left, top, right,
bottom) for the bounds of the source rectangle

Performance Note: Not currently implemented in-place in the core layer.

close(self)
Closes the file pointer, if possible.

This operation will destroy the image core and release its memory.
The image data will be unusable afterward.

This function is only required to close images that have not
had their file read and closed by the :py:meth:`~PIL.Image.Image.load` method. See :ref:`file-handling` for more information.

convert(self, mode=None, matrix=None, dither=None, palette=0, colors=256)
Returns a converted copy of this image. For the "P" mode, this
method translates pixels through the palette. If mode is omitted, a mode is chosen so that all information in the image
and the palette can be represented without a palette.

The current version supports all possible conversions between
ports "L", "RGB" and "CMYK." The ****matrix**** argument only supports "L" and "RGB".

When translating a color image to greyscale (mode "L"), the library uses the ITU-R 601-2 luma transform::

$$L = R * 299/1000 + G * 587/1000 + B * 114/1000$$

The default method of converting a greyscale ("L") or "RGB" image into a bilevel (mode "1") image uses Floyd-Steinberg dither to approximate the original image luminosity levels. If dither is NONE, all values larger than 128 are set to 255 (white), all other values to 0 (black). To use other thresholds, use the :py:meth:`~PIL.Image.Image.point` method.

When converting from "RGBA" to "P" without a ****matrix**** argument, this passes the operation to :py:meth:`~PIL.Image.Image.quantize`,

and `**dither**` and `**palette**` are ignored.

`:param mode:` The requested mode. See: `:ref:`concept-mode``.

`:param matrix:` An optional conversion matrix. If given, should be 4- or 12-tuple containing floating point values.

`:param dither:` Dithering method, used when converting from mode "RGB" to "P" or from "RGB" or "L" to "1". Available methods are NONE or FLOYDSTEINBERG (default).

Note that this is not used when `**matrix**` is supplied.

`:param palette:` Palette to use when converting from mode "RGB" to "P". Available palettes are WEB or ADAPTIVE.

`:param colors:` Number of colors to use for the ADAPTIVE palette. Defaults to 256.

`:rtype:` `:py:class:`~PIL.Image.Image``

`:returns:` An `:py:class:`~PIL.Image.Image`` object.

`copy(self)`
Copies this image. Use this method if you wish to paste into an image, but still retain the original.

`:rtype:` `:py:class:`~PIL.Image.Image``

`:returns:` An `:py:class:`~PIL.Image.Image`` object.

`crop(self, box=None)`
Returns a rectangular region from this image. The box is a 4-tuple defining the left, upper, right, and lower pixel coordinate. See `:ref:`coordinate-system``.

Note: Prior to Pillow 3.4.0, this was a lazy operation.

`:param box:` The crop rectangle, as a (left, upper, right, lower)-tuple.

`:rtype:` `:py:class:`~PIL.Image.Image``

`:returns:` An `:py:class:`~PIL.Image.Image`` object.

`draft(self, mode, size)`
Configures the image file loader so it returns a version of the image that as closely as possible matches the given mode and size. For example, you can use this method to convert a color JPEG to greyscale while loading it, or to extract a 128x192 version from a PCD file.

Note that this method modifies the `:py:class:`~PIL.Image`` object in place. If the image has already been loaded, this method has no effect.

is

Note: This method is not implemented for most images. It currently implemented only for JPEG and PCD images.

:param mode: The requested mode.

:param size: The requested size.

effect_spread(self, distance)

Randomly spread pixels in an image.

:param distance: Distance to spread pixels.

filter(self, filter)

Filters this image using the given filter. For a list o

f

available filters, see the :py:mod:`~PIL.ImageFilter` mo

dule.

:param filter: Filter kernel.

:returns: An :py:class:`~PIL.Image.Image` object.

frombytes(self, data, decoder_name='raw', *args)

Loads this image with pixel data from a bytes object.

This method is similar to the :py:func:`~PIL.Image.fromb
ytes` function,

but loads data into this image instead of creating a new
image object.

fromstring(self, *args, **kw)

getbands(self)

Returns a tuple containing the name of each band in this

image.

For example, `**getbands**` on an RGB image returns ("R",
"G", "B").

"G", "B").

:returns: A tuple containing band names.

:rtype: tuple

getbbox(self)

Calculates the bounding box of the non-zero regions in t

he

image.

:returns: The bounding box is returned as a 4-tuple defi

ning the

left, upper, right, and lower pixel coordinate. See

:ref:`coordinate-system`. If the image is completely

empty, this

method returns None.

getchannel(self, channel)

Returns an image containing a single channel of the sour

ce image.

:param channel: What channel to return. Could be index

(0 for "R" channel of "RGB") or channel name

("A" for alpha channel of "RGBA").

:returns: An image in "L" mode.

```

.. versionadded:: 4.3.0

getcolors(self, maxcolors=256)
    Returns a list of colors used in this image.

    :param maxcolors: Maximum number of colors. If this num
ber is exceeded, this method returns None. The default limi
t is 256 colors.
    :returns: An unsorted list of (count, pixel) values.

getdata(self, band=None)
    Returns the contents of this image as a sequence object
    containing pixel values. The sequence object is flatten
ed, so that values for line one follow directly after the value
s of line zero, and so on.

    Note that the sequence object returned by this method is
an internal PIL data type, which only supports certain sequ
ence operations. To convert it to an ordinary sequence (e.g.
for printing), use **list(im.getdata())**.

    :param band: What band to return. The default is to ret
urn all bands. To return a single band, pass in the inde
x value (e.g. 0 to get the "R" band from an "RGB" imag
e).
    :returns: A sequence-like object.

getextrema(self)
    Gets the the minimum and maximum pixel values for each b
and in the image.

    :returns: For a single-band image, a 2-tuple containing
the minimum and maximum pixel value. For a multi-band im
age, a tuple containing one 2-tuple for each band.

getim(self)
    Returns a capsule that points to the internal image memo
ry.

    :returns: A capsule object.

getpalette(self)
    Returns the image palette as a list.

    :returns: A list of color values [r, g, b, ...], or None
if the image has no palette.

```

```

|   getpixel(self, xy)
|       Returns the pixel value at a given position.
|
|       :param xy: The coordinate, given as (x, y). See
|           :ref:`coordinate-system`.
|       :returns: The pixel value. If the image is a multi-laye
r image,
|           this method returns a tuple.
|
|   getprojection(self)
|       Get projection to x and y axes
|
|       :returns: Two sequences, indicating where there are non-
zero
|           pixels along the X-axis and the Y-axis, respectivel
y.
|
|   histogram(self, mask=None, extrema=None)
|       Returns a histogram for the image. The histogram is retu
rned as
|       a list of pixel counts, one for each pixel value in the
source
|       image. If the image has more than one band, the histogra
ms for
|       all bands are concatenated (for example, the histogram f
or an
|       "RGB" image contains 768 values).
|
|       A bilevel image (mode "1") is treated as a greyscale
("L") image
|       by this method.
|
|       If a mask is provided, the method returns a histogram fo
r those
|       parts of the image where the mask image is non-zero. The
mask
|       image must have the same size as the image, and be eithe
r a
|       bi-level image (mode "1") or a greyscale image ("L").
|
|       :param mask: An optional mask.
|       :returns: A list containing pixel counts.
|
|   load(self)
|       Allocates storage for the image and loads the pixel dat
a. In
|       normal cases, you don't need to call this method, since
the
|       Image class automatically loads an opened image when it
is
|       accessed for the first time.
|
|       If the file associated with the image was opened by Pill
ow, then this
|       method will close it. The exception to this is if the im
age has
|       multiple frames, in which case the file will be left ope
n for seek
|       operations. See :ref:`file-handling` for more informatio
n.

```

```

:returns: An image access object.
:rtype: :ref:`PixelAccess` or :py:class:`PIL.PyAccess`

offset(self, xoffset, yoffset=None)

paste(self, im, box=None, mask=None)
    Pastes another image into this image. The box argument i
s either
    a 2-tuple giving the upper left corner, a 4-tuple defini
ng the
    left, upper, right, and lower pixel coordinate, or None
(same as
    (0, 0)). See :ref:`coordinate-system`. If a 4-tuple is g
iven, the size
    of the pasted image must match the size of the region.

    If the modes don't match, the pasted image is converted
to the mode of
this image (see the :py:meth:`~PIL.Image.Image.convert`
method for
details).

    Instead of an image, the source can be a integer or tupl
e
    containing pixel values. The method then fills the regi
on
    with the given color. When creating RGB images, you can
also use color strings as supported by the ImageColor mo
dule.

    If a mask is given, this method updates only the regions
indicated by the mask. You can use either "1", "L" or
"RGBA"
images (in the latter case, the alpha band is used as ma
sk).
Where
the mask is 255, the given image is copied as is.
iate
the mask is 0, the current value is preserved. Intermed
alpha
values will mix the two images together, including their
channels if they have them.

See :py:meth:`~PIL.Image.Image.alpha_composite` if you w
ant to
combine images with respect to their alpha channels.

:param im: Source image or pixel value (integer or tupl
e).
te into.
    :param box: An optional 4-tuple giving the region to pas
te into.
        If a 2-tuple is used instead, it's treated as the upp
er left
        corner. If omitted or None, the source is pasted int
o the
        upper left corner.

        If an image is given as the second argument and there
is no
third, the box defaults to (0, 0), and the second arg
ument

```


is interpreted as a mask image.
:param mask: An optional mask image.

point(self, lut, mode=None)

Maps this image through a lookup table or function.

:param lut: A lookup table, containing 256 (or 65536 if
self.mode=="I" and mode == "L") values per band in the

image. A function can be used instead, it should take

single argument. The function is called once for each
possible pixel value, and the resulting table is applied

all bands of the image.

:param mode: Output mode (default is same as input). In

current version, this can only be used if the source

has mode "L" or "P", and the output has mode "l" or the

source image mode is "I" and the output mode is "L".

:returns: An :py:class:`~PIL.Image.Image` object.

putalpha(self, alpha)

Adds or replaces the alpha layer in this image. If the

does not have an alpha layer, it's converted to "LA" or

"RGBA". The new layer must be either "L" or "l".

:param alpha: The new alpha layer. This can either be a

image having the same size as this image, or an integer or

other color value.

putdata(self, data, scale=1.0, offset=0.0)

Copies pixel data to this image. This method copies data

from a sequence object into the image, starting at the upper left

corner (0, 0), and continuing until either the image or

sequence ends. The scale and offset values are used to

adjust the sequence values: `**pixel = value*scale + offset**`.

:param data: A sequence object.

:param scale: An optional scale value. The default is

1.0. :param offset: An optional offset value. The default is

0.0. putpalette(self, data, rawmode='RGB')

Attaches a palette to this image. The image must be a

"P" or "L" image, and the palette sequence must contain 768 integer

values, where each group of three values represent the red,

green, and blue values for the corresponding pixel index. Instead of an integer sequence, you can use an 8-bit string.

`:param data:` A palette sequence (either a list or a string).

`:param rawmode:` The raw mode of the palette.

`putpixel(self, xy, value)`
 Modifies the pixel at the given position. The color is given as a single numerical value for single-band images, and a tuple for multi-band images. In addition to this, RGB and RGBA tuples are accepted for P images.

Note that this method is relatively slow. For more extensive changes, use `:py:meth:`~PIL.Image.Image.paste`` or the `:py:mod:`~PIL.ImageDraw`` module instead.

See:

- * `:py:meth:`~PIL.Image.Image.paste``
- * `:py:meth:`~PIL.Image.Image.putdata``
- * `:py:mod:`~PIL.ImageDraw``

`:param xy:` The pixel coordinate, given as (x, y). See `:ref:`coordinate-system``.

`:param value:` The pixel value.

`quantize(self, colors=256, method=None, kmeans=0, palette=None)`
 Convert the image to 'P' mode with the specified number of colors.

`:param colors:` The desired number of colors, <= 256

`:param method:` 0 = median cut
 1 = maximum coverage
 2 = fast octree
 3 = libimagequant

`:param kmeans:` Integer

`:param palette:` Quantize to the palette of given `:py:class:`~PIL.Image.Image``.

`:returns:` A new image

`remap_palette(self, dest_map, source_palette=None)`
 Rewrites the image to reorder the palette.

`:param dest_map:` A list of indexes into the original palette.
 e.g. [1,0] would swap a two item palette, and `list(range(255))` is the identity transform.

`:param source_palette:` Bytes or None.

`:returns:` An `:py:class:`~PIL.Image.Image`` object.

`resize(self, size, resample=0, box=None)`

```

|         Returns a resized copy of this image.
|
|         :param size: The requested size in pixels, as a 2-tuple:
|             (width, height).
|         :param resample: An optional resampling filter. This can
n be
|             one of :py:attr:`PIL.Image.NEAREST`, :py:attr:`PIL.Im
age.BOX`,
|             :py:attr:`PIL.Image.BILINEAR`, :py:attr:`PIL.Image.HA
MMING`,
|             :py:attr:`PIL.Image.BICUBIC` or :py:attr:`PIL.Image.L
ANCZOS`.
|             If omitted, or if the image has mode "1" or "P", it is
s
|             set :py:attr:`PIL.Image.NEAREST`.
|             See: :ref:`concept-filters`.
|         :param box: An optional 4-tuple of floats giving the region
ion
|             of the source image which should be scaled.
|             The values should be within (0, 0, width, height) rectangle.
|             If omitted or None, the entire source is used.
|         :returns: An :py:class:`~PIL.Image.Image` object.
|
|         rotate(self, angle, resample=0, expand=0, center=None, translate=None, fillcolor=None)
|         Returns a rotated copy of this image. This method returns a
ns a
|         copy of this image, rotated the given number of degrees
counter
|         clockwise around its centre.
|
|         :param angle: In degrees counter clockwise.
|         :param resample: An optional resampling filter. This can
n be
|             one of :py:attr:`PIL.Image.NEAREST` (use nearest neighbour),
|             :py:attr:`PIL.Image.BILINEAR` (linear interpolation in a 2x2
|             environment), or :py:attr:`PIL.Image.BICUBIC` (cubic spline interpolation in a 4x4 environment).
|             If omitted, or if the image has mode "1" or "P", it is
s
|             set :py:attr:`PIL.Image.NEAREST`. See :ref:`concept-filters`.
|         :param expand: Optional expansion flag. If true, expands the output
s the output
|         image to make it large enough to hold the entire rotated image.
|         If false or omitted, make the output image the same size as the
|         input image. Note that the expand flag assumes rotation around
ion around
|         the center and no translation.
|         :param center: Optional center of rotation (a 2-tuple).
|             Origin is
|             the upper left corner. Default is the center of the
image.
|         :param translate: An optional post-rotate translation (a
2-tuple).
|         :param fillcolor: An optional color for area outside the

```

rotated image.

:returns: An :py:class:`~PIL.Image.Image` object.

save(self, fp, format=None, **params)

Saves this image under the given filename. If no format

is

specified, the format to use is determined from the file

name

extension, if possible.

Keyword options can be used to provide additional instru

ctions

to the writer. If a writer doesn't recognise an option,

it is

silently ignored. The available options are described in

the

:doc:`image format documentation

<../handbook/image-file-formats>` for each writer.

You can use a file object instead of a filename. In this

case,

you must always specify the format. The file object must implement the ``seek``, ``tell``, and ``write`` methods, and be opened in binary mode.

:param fp: A filename (string), pathlib.Path object or f

ile object.

:param format: Optional format override. If omitted, th

e

format to use is determined from the filename extensi

on.

If a file object was used instead of a filename, this parameter should always be used.

:param params: Extra parameters to the image writer.

:returns: None

:exception ValueError: If the output format could not be

determined

from the file name. Use the format option to solve t

his.

:exception IOError: If the file could not be written. T

he file

may have been created, and may contain partial data.

seek(self, frame)

Seeks to the given frame in this sequence file. If you s

seek

beyond the end of the sequence, the method raises an **EOFError** exception. When a sequence file is opened,

the

library automatically seeks to frame 0.

Note that in the current version of the library, most se

quence

formats only allows you to seek to the next frame.

See :py:meth:`~PIL.Image.Image.tell`.

:param frame: Frame number, starting at 0.

:exception EOFError: If the call attempts to seek beyond

the end

of the sequence.

```

show(self, title=None, command=None)
    Displays this image. This method is mainly intended for
    debugging purposes.

    On Unix platforms, this method saves the image to a temp
orary
PPM file, and calls either the **xv** utility or the **d
isplay**
utility, depending on which one can be found.

    On macOS, this method saves the image to a temporary BMP
file, and
opens it with the native Preview application.

    On Windows, it saves the image to a temporary BMP file,
and uses
the standard BMP display utility to show it (usually Pai
nt).

    :param title: Optional title to use for the image windo
w,
    where possible.
    :param command: command used to show the image

split(self)
    Split this image into individual bands. This method retu
rns a
tuple of individual image bands from an image. For examp
le,
splitting an "RGB" image creates three new images each
en,
containing a copy of one of the original bands (red, gre
en,
blue).

    If you need only one band, :py:meth:`~PIL.Image.Image.ge
tchannel`
method can be more convenient and faster.

    :returns: A tuple containing bands.

tell(self)
    Returns the current frame number. See :py:meth:`~PIL.Ima
ge.Image.seek`.

    :returns: Frame number, starting with 0.

thumbnail(self, size, resample=3)
    Make this image into a thumbnail. This method modifies
the
image to contain a thumbnail version of itself, no large
r than
the given size. This method calculates an appropriate t
humbnail
size to preserve the aspect of the image, calls the
:py:meth:`~PIL.Image.Image.draft` method to configure th
e file reader
(where applicable), and finally resizes the image.

    Note that this function modifies the :py:class:`~PIL.Ima
ge.Image`

```

```

        object in place. If you need to use the full resolution
image as well,
        apply this method to a :py:meth:`~PIL.Image.Image.copy`
of the original
        image.

        :param size: Requested size.
        :param resample: Optional resampling filter. This can b
e one
        of :py:attr:`~PIL.Image.NEAREST`, :py:attr:`~PIL.Image.
BILINEAR`,
        :py:attr:`~PIL.Image.BICUBIC`, or :py:attr:`~PIL.Image.
LANCZOS`.
        If omitted, it defaults to :py:attr:`~PIL.Image.BICUBI
C`.
        (was :py:attr:`~PIL.Image.NEAREST` prior to version 2.
5.0)

        :returns: None

tobitmap(self, name='image')
    Returns the image converted to an X11 bitmap.

    .. note:: This method only works for mode "1" images.

    :param name: The name prefix to use for the bitmap varia
bles.

    :returns: A string containing an X11 bitmap.
    :raises ValueError: If the mode is not "1"

tobytes(self, encoder_name='raw', *args)
    Return image as a bytes object.

    .. warning::

        This method returns the raw image data from the inte
rnal
        storage. For compressed image data (e.g. PNG, JPEG)
use
        :meth:`~.save`, with a BytesIO parameter for in-memo
ry
        data.

        :param encoder_name: What encoder to use. The default i
s to
        use the standard "raw" encoder.
        :param args: Extra arguments to the encoder.
        :rtype: A bytes object.

toqimage(self)
    Returns a QImage copy of this image

toqixmap(self)
    Returns a QPixmap copy of this image

tostring(self, *args, **kw)

transform(self, size, method, data=None, resample=0, fill=1,
fillcolor=None)
    Transforms this image. This method creates a new image
with the
    given size, and the same mode as the original, and copie

```

```

s data
    to the new image using the given transform.

    :param size: The output size.
    :param method: The transformation method. This is one o
f
    :py:attr:`PIL.Image.EXTENT` (cut out a rectangular sub
region),
    :py:attr:`PIL.Image.AFFINE` (affine transform),
    :py:attr:`PIL.Image.PERSPECTIVE` (perspective transfor
m),
    :py:attr:`PIL.Image.QUAD` (map a quadrilateral to a re
ctangle), or
    :py:attr:`PIL.Image.MESH` (map a number of source quad
rilaterals
    in one operation).

    It may also be an :py:class:`~PIL.Image.ImageTransform
Handler`
    object::
        class Example(Image.ImageTransformHandler):
            def transform(size, method, data, resample, fill
=1):
                # Return result

    It may also be an object with a :py:meth:`~method.getda
ta` method
    that returns a tuple supplying new **method** and **da
ta** values::
        class Example(object):
            def getdata(self):
                method = Image.EXTENT
                data = (0, 0, 100, 100)
                return method, data
    :param data: Extra data to the transformation method.
    :param resample: Optional resampling filter. It can be
one of
    :py:attr:`PIL.Image.NEAREST` (use nearest neighbour),
    :py:attr:`PIL.Image.BILINEAR` (linear interpolation in
n a 2x2
    environment), or :py:attr:`PIL.Image.BICUBIC` (cubic
spline
    interpolation in a 4x4 environment). If omitted, or i
f the image
    has mode "1" or "P", it is set to :py:attr:`PIL.Imag
e.NEAREST`.
    :param fill: If **method** is an
    :py:class:`~PIL.Image.ImageTransformHandler` object, t
his is one of
    the arguments passed to it. Otherwise, it is unused.
    :param fillcolor: Optional fill color for the area outsi
de the
    transform in the output image.
    :returns: An :py:class:`~PIL.Image.Image` object.

    transpose(self, method)
    Transpose image (flip or rotate in 90 degree steps)

    :param method: One of :py:attr:`PIL.Image.FLIP_LEFT_RIGH
T`,
    :py:attr:`PIL.Image.FLIP_TOP_BOTTOM`, :py:attr:`PIL.Im

```

```

age.ROTATE_90`,
|           :py:attr:`PIL.Image.ROTATE_180`, :py:attr:`PIL.Image.R
OTATE_270`,
|           :py:attr:`PIL.Image.TRANSPOSE` or :py:attr:`PIL.Image.
TRANSVERSE`.
|           :returns: Returns a flipped or rotated copy of this imag
e.

```

```

|         verify(self)
|             Verifies the contents of a file. For data read from a fi
le, this
|             method attempts to determine if the file is broken, with
out
|             actually decoding the image data. If this method finds
any
|             problems, it raises suitable exceptions. If you need to
load
|             the image after using this method, you must reopen the i
mage
|             file.

```

Data descriptors defined here:

```

|     __array_interface__
|
|     __dict__
|         dictionary for instance variables (if defined)
|
|     __weakref__
|         list of weak references to the object (if defined)
|
|     height
|
|     size
|
|     width

```

Data and other attributes defined here:

```

|     __hash__ = None
|
|     format = None
|
|     format_description = None

```

```

class ImagePointHandler(builtins.object)

```

Data descriptors defined here:

```

|     __dict__
|         dictionary for instance variables (if defined)
|
|     __weakref__
|         list of weak references to the object (if defined)

```

```

class ImageTransformHandler(builtins.object)

```

Data descriptors defined here:


```

|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)

```

FUNCTIONS

```

alpha_composite(im1, im2)
    Alpha composite im2 over im1.

```

```

:param im1: The first image. Must have mode RGBA.
:param im2: The second image. Must have mode RGBA, and the
same size as
    the first image.
:returns: An :py:class:`~PIL.Image.Image` object.

```

```

blend(im1, im2, alpha)
    Creates a new image by interpolating between two input image
s, using
    a constant alpha.:

```

```

    out = image1 * (1.0 - alpha) + image2 * alpha

```

```

:param im1: The first image.
:param im2: The second image. Must have the same mode and s
ize as
    the first image.
:param alpha: The interpolation alpha factor. If alpha is
0.0, a
    copy of the first image is returned. If alpha is 1.0, a c
opy of
    the second image is returned. There are no restrictions o
n the
    alpha value. If necessary, the result is clipped to fit i
nto
    the allowed output range.
:returns: An :py:class:`~PIL.Image.Image` object.

```

```

coerce_e(value)

```

```

composite(image1, image2, mask)
    Create composite image by blending images using a transparen
cy mask.

```

```

:param image1: The first image.
:param image2: The second image. Must have the same mode an
d
    size as the first image.
:param mask: A mask image. This image can have mode
    "1", "L", or "RGBA", and must have the same size as the
    other two images.

```

```

effect_mandelbrot(size, extent, quality)
    Generate a Mandelbrot set covering the given extent.

```

```

:param size: The requested size in pixels, as a 2-tuple:
    (width, height).
:param extent: The extent to cover, as a 4-tuple:
    (x0, y0, x1, y2).
:param quality: Quality.

```

```

effect_noise(size, sigma)
    Generate Gaussian noise centered around 128.

:param size: The requested size in pixels, as a 2-tuple:
    (width, height).
:param sigma: Standard deviation of noise.

eval(image, *args)
    Applies the function (which should take one argument) to each
    pixel in the given image. If the image has more than one band, the
    same function is applied to each band. Note that the function is
    evaluated once for each possible pixel value, so you cannot
    use random components or other generators.

:param image: The input image.
:param function: A function object, taking one integer argument.
:returns: An :py:class:`~PIL.Image.Image` object.

fromarray(obj, mode=None)
    Creates an image memory from an object exporting the array interface
    (using the buffer protocol).

    If obj is not contiguous, then the tobytes method is called
    and :py:func:`~PIL.Image.frombuffer` is used.

    If you have an image in NumPy::

        from PIL import Image
        import numpy as np
        im = Image.open('hopper.jpg')
        a = np.asarray(im)

    Then this can be used to convert it to a Pillow image::

        im = Image.fromarray(a)

:param obj: Object with array interface
:param mode: Mode to use (will be determined from type if None)
    See :ref:`concept-modes`.
:returns: An image object.

.. versionadded:: 1.1.6

frombuffer(mode, size, data, decoder_name='raw', *args)
    Creates an image memory referencing pixel data in a byte buffer.

    This function is similar to :py:func:`~PIL.Image.frombytes`,
    but uses data in the byte buffer, where possible. This means that changes
    to the original buffer object are reflected in this image). Not all
    modes can share memory; supported modes include "L", "RGBX", "RGBA", a

```

nd "CMYK".

Note that this function decodes pixel data only, not entire images.

If you have an entire image file in a string, wrap it in a `**BytesIO**` object, and use `:py:func:`~PIL.Image.open`` to load it.

In the current version, the default parameters used for the "raw" decoder

differs from that used for `:py:func:`~PIL.Image.frombytes``. This is a

bug, and will probably be fixed in a future release. The current release

issues a warning if you do this; to disable the warning, you should provide

the full set of parameters. See below for details.

`:param mode:` The image mode. See: `:ref:`~concept-modes``.

`:param size:` The image size.

`:param data:` A bytes or other buffer object containing raw data for the given mode.

`:param decoder_name:` What decoder to use.

`:param args:` Additional parameters for the given decoder. For

or the

default encoder ("raw"), it's recommended that you provide

de the

full set of parameters::

```
frombuffer(mode, size, data, "raw", mode, 0, 1)
```

`:returns:` An `:py:class:`~PIL.Image.Image`` object.

`.. versionadded:: 1.1.4`

```
frombytes(mode, size, data, decoder_name='raw', *args)
```

Creates a copy of an image memory from pixel data in a buffer.

In its simplest form, this function takes three arguments (mode, size, and unpacked pixel data).

You can also use any pixel decoder supported by PIL. For more

information on available decoders, see the section `:ref:`~Writing Your Own File Decoder <file-decoders>``.

Note that this function decodes pixel data only, not entire images.

If you have an entire image in a string, wrap it in a `:py:class:`~io.BytesIO`` object, and use `:py:func:`~PIL.Image.open`` to load it.

`:param mode:` The image mode. See: `:ref:`~concept-modes``.

`:param size:` The image size.

`:param data:` A byte buffer containing raw data for the given

mode.

`:param decoder_name:` What decoder to use.

`:param args:` Additional parameters for the given decoder.

`:returns:` An `:py:class:`~PIL.Image.Image`` object.

```

fromqimage(im)
    Creates an image instance from a QImage image

fromqixmap(im)
    Creates an image instance from a QPixmap image

fromstring(*args, **kw)

getmodebandnames(mode)
    Gets a list of individual band names. Given a mode, this fu
nction returns
    a tuple containing the names of individual bands (use
    :py:method:`~PIL.Image.getmodetype` to get the mode used to
store each
    individual band.

    :param mode: Input mode.
    :returns: A tuple containing band names. The length of the
tuple
        gives the number of bands in an image of the given mode.
    :exception KeyError: If the input mode was not a standard mo
de.

getmodebands(mode)
    Gets the number of individual bands for this mode.

    :param mode: Input mode.
    :returns: The number of bands in this mode.
    :exception KeyError: If the input mode was not a standard mo
de.

getmodebase(mode)
    Gets the "base" mode for given mode. This function returns
"L" for
    images that contain grayscale data, and "RGB" for images tha
t
    contain color data.

    :param mode: Input mode.
    :returns: "L" or "RGB".
    :exception KeyError: If the input mode was not a standard mo
de.

getmodetype(mode)
    Gets the storage type mode. Given a mode, this function ret
urns a
    single-layer mode suitable for storing individual bands.

    :param mode: Input mode.
    :returns: "L", "I", or "F".
    :exception KeyError: If the input mode was not a standard mo
de.

init()
    Explicitly initializes the Python Imaging Library. This func
tion
    loads all available file format drivers.

isImageType(t)
    Checks if an object is an image object.

```

.. warning::

This function is for internal use only.

:param t: object to check if it's an image

:returns: True if the object is an image

linear_gradient(mode)

Generate 256x256 linear gradient from black to white, top to bottom.

:param mode: Input mode.

merge(mode, bands)

Merge a set of single band images into a new multiband image.

:param mode: The mode to use for the output image. See:

:ref:`concept-modes`.

:param bands: A sequence containing one single-band image for

each band in the output image. All bands must have the same size.

:returns: An :py:class:`~PIL.Image.Image` object.

new(mode, size, color=0)

Creates a new image with the given mode and size.

:param mode: The mode to use for the new image. See:

:ref:`concept-modes`.

:param size: A 2-tuple, containing (width, height) in pixels.

:param color: What color to use for the image. Default is black.

If given, this should be a single integer or floating point value

for single-band modes, and a tuple for multi-band modes

(one value

per band). When creating RGB images, you can also use color

strings as supported by the ImageColor module. If the color is

None, the image is not initialised.

:returns: An :py:class:`~PIL.Image.Image` object.

open(fp, mode='r')

Opens and identifies the given image file.

This is a lazy operation; this function identifies the file, but

the file remains open and the actual image data is not read from

the file until you try to process the data (or call the

:py:meth:`~PIL.Image.Image.load` method). See

:py:func:`~PIL.Image.new`. See :ref:`file-handling`.

:param fp: A filename (string), pathlib.Path object or a file object.

The file object must implement :py:meth:`~file.read`,

:py:meth:`~file.seek`, and :py:meth:`~file.tell` methods,

and be opened in binary mode.
:param mode: The mode. If given, this argument must be "r".
:returns: An :py:class:`~PIL.Image.Image` object.
:exception IOError: If the file cannot be found, or the image cannot be opened and identified.

preinit()
Explicitly load standard file format drivers.

radial_gradient(mode)
Generate 256x256 radial gradient from black to white, centre to edge.

:param mode: Input mode.

register_decoder(name, decoder)
Registers an image decoder. This function should not be used in application code.

:param name: The name of the decoder
:param decoder: A callable(mode, args) that returns an ImageFile.PyDecoder object

.. versionadded:: 4.1.0

register_encoder(name, encoder)
Registers an image encoder. This function should not be used in application code.

:param name: The name of the encoder
:param encoder: A callable(mode, args) that returns an ImageFile.PyEncoder object

.. versionadded:: 4.1.0

register_extension(id, extension)
Registers an image extension. This function should not be used in application code.

:param id: An image format identifier.
:param extension: An extension used for this format.

register_extensions(id, extensions)
Registers image extensions. This function should not be used in application code.

:param id: An image format identifier.
:param extensions: A list of extensions used for this format.

t.

register_mime(id, mimetype)
Registers an image MIME type. This function should not be used in application code.

:param id: An image format identifier.
:param mimetype: The image MIME type for this format.

register_open(id, factory, accept=None)
Register an image file plugin. This function should not be

used

in application code.

:param id: An image format identifier.

:param factory: An image file factory method.

:param accept: An optional function that can be used to quick

ly

reject images having another format.

register_save(id, driver)

Registers an image save function. This function should not

be

used in application code.

:param id: An image format identifier.

:param driver: A function to save images in this format.

register_save_all(id, driver)

Registers an image function to save all the frames of a multiframe format. This function should not be used in application code.

:param id: An image format identifier.

:param driver: A function to save images in this format.

registered_extensions()

Returns a dictionary containing all file extensions belonging

g

to registered plugins

DATA

ADAPTIVE = 1

AFFINE = 0

ANTIALIAS = 1

BICUBIC = 3

BILINEAR = 2

BOX = 4

CONTAINER = 2

CUBIC = 3

DECODERS = {}

DEFAULT_STRATEGY = 0

ENCODERS = {}

EXTENSION = {}

EXTENT = 1

FASTOCTREE = 2

FILTERED = 1

FIXED = 4

FLIP_LEFT_RIGHT = 0

FLIP_TOP_BOTTOM = 1

FLOYDSTEINBERG = 3

HAMMING = 5

HAS_PATHLIB = True

HUFFMAN_ONLY = 2

ID = []

LANCZOS = 1

LIBIMAGEQUANT = 3

LINEAR = 2

MAXCOVERAGE = 1

MAX_IMAGE_PIXELS = 89478485

MEDIANCUT = 0

MESH = 4

```
MIME = {}
MODES = ['1', 'CMYK', 'F', 'HSV', 'I', 'L', 'LAB', 'P', 'RGB',
'RGBA',...
NEAREST = 0
NONE = 0
NORMAL = 0
OPEN = {}
ORDERED = 1
PERSPECTIVE = 2
PILLOW_VERSION = '5.4.1'
QUAD = 3
RASTERIZE = 2
RLE = 3
ROTATE_180 = 3
ROTATE_270 = 4
ROTATE_90 = 2
SAVE = {}
SAVE_ALL = {}
SEQUENCE = 1
TRANSPPOSE = 5
TRANSVERSE = 6
USE_CFFI_ACCESS = False
VERSION = '1.1.7'
WEB = 0
logger = <Logger PIL.Image (WARNING)>
py3 = True
```

```
VERSION
5.4.1
```

```
FILE
/opt/conda/lib/python3.7/site-packages/PIL/Image.py
```

Running `help()` on `Image` tells us that this object is "the Image class wrapper". We see from the top level documentation about the image object that there is "hardly ever any reason to call the Image constructor directly", and they suggest that the open function might be the way to go.

In [9]:

```
# Lets call help on the open function to see what it's all about. Remember that  
since we want to pass in the  
# function reference, and not run the function itself, we don't put parentheses b  
ehind the function name.  
help(Image.open)
```

Help on function open in module PIL.Image:

```
open(fp, mode='r')  
    Opens and identifies the given image file.  
  
    This is a lazy operation; this function identifies the file, but  
    the file remains open and the actual image data is not read from  
    the file until you try to process the data (or call the  
    :py:meth:`~PIL.Image.Image.load` method). See  
    :py:func:`~PIL.Image.new`. See :ref:`file-handling`.  
  
    :param fp: A filename (string), pathlib.Path object or a file ob  
ject.  
        The file object must implement :py:meth:`~file.read`,  
        :py:meth:`~file.seek`, and :py:meth:`~file.tell` methods,  
        and be opened in binary mode.  
    :param mode: The mode. If given, this argument must be "r".  
    :returns: An :py:class:`~PIL.Image.Image` object.  
    :exception IOError: If the file cannot be found, or the image ca  
nnot be  
        opened and identified.
```

In [10]:

```
# It looks like Image.open() is a function that loads an image from a file and r  
eturns an instance  
# of the Image class. Lets give it a try. In the read_only directory there is an  
image I've provided  
# which is from our Master's of Information program recruitment flyer. Lets try  
and load that now
```

```
file="readonly/msi_recruitment.gif"  
image=Image.open(file)  
print(image)
```

```
<PIL.GifImagePlugin.GifImageFile image mode=P size=800x450 at 0x7FEC  
344F4EB8>
```

getmro() to examine Child Classes

In [12]:

```
# Ok, we see that this returns us a kind of PIL.GifImagePlugin.GifImageFile. At first this might seem a bit confusing, since because we were told by the docs that we should be expecting a PIL.Image.Image object back. But this is just object inheritance working! In fact, the object returned is both an Image and a GifImageFile. We can use the python inspect module to see this
# as the getmro function will return a list of all of the classes that are being inherited by a given object. Lets try it.

import inspect
print("The type of the image is " + str(type(image)))
inspect.getmro(type(image))
# Note that the Image.Image is the most general, (the root)
#and GifImagePlugin.GifImageFile is the most specific (lowest on the inheritance chain)
```

The type of the image is <class 'PIL.GifImagePlugin.GifImageFile'>

Out[12]:

```
(PIL.GifImagePlugin.GifImageFile,
PIL.ImageFile.ImageFile,
PIL.Image.Image,
object)
```

Showing Images on `.show()`

In [13]:

```
# Now that we are comfortable with the object. How do we view the image? It turns out that the image object has a show function. You can find this by looking at all of the properties of the object if you wanted to, using the dir() function.
image.show()
```

Using `IPython.display` to show images

In [14]:

```
# Hrm, that didn't seem to have the intended effect. The problem is that the image is stored
# remotely, on Coursera's server, but show tries to show it locally to you. So, if the Coursera
# server software was running on someone's workstation in Mountain View California, where Coursera
# has its offices, then you just popped up a picture of our recruitment materials. Thanks! :)
# Instead, we want to render the image in the Jupyter notebook. It turns out Jupyter has a function
# which can help with this.
from IPython.display import display
display(image)
```



For those who would like to understand this in more detail, the Jupyter environment is running a **special wrapper around the Python interpreter, called IPython**. IPython allows the kernel back end to communicate with a browser front end, among other things. The IPython package has a display function which can take objects and use custom formatters in order to render them. A number of formatters are provided by default, including one which knows how to handle image types.

That's a quick overview of how to read and display images using pillow, in the next lecture we'll jump in a bit more detail to understand how to use pillow to manipulate images.

Common Functions in the Python Imaging Library

Lets take a look at some of the common tasks we can do in python using the pillow library.

In [15]:

```
# First, lets import the PIL library and the Image object
import PIL
from PIL import Image
# And lets import the display functionality
from IPython.display import display
# And finally, lets load the image we were working with last time
file="readonly/msi_recruitment.gif"
image=Image.open(file)
```

Using `image.copy` and `image.save`

In [16]:

```
# Great, now lets check out a few more methods of the image library. First, we'll
look at copy
# And if you remember, we can do this using the built in python help() function
help(image.copy)
```

Help on method copy in module PIL.Image:

`copy()` method of `PIL.GifImagePlugin.GifImageFile` instance
Copies this image. Use this method if you wish to paste things
into an image, but still retain the original.

`:rtype:` `:py:class:`~PIL.Image.Image``
`:returns:` An `:py:class:`~PIL.Image.Image`` object.

In [17]:

```
# We can see that copy takes no arguments, and that the return object is an Image object itself. Now lets  
# look at save  
help(image.save)
```

Help on method save in module PIL.Image:

save(fp, format=None, **params) method of PIL.GifImagePlugin.GifImageFile instance

Saves this image under the given filename. If no format is specified, the format to use is determined from the filename extension, if possible.

Keyword options can be used to provide additional instructions to the writer. If a writer doesn't recognise an option, it is silently ignored. The available options are described in the :doc:`image format documentation <../handbook/image-file-formats>` for each writer.

You can use a file object instead of a filename. In this case, you must always specify the format. The file object must implement the ``seek``, ``tell``, and ``write`` methods, and be opened in binary mode.

:param fp: A filename (string), pathlib.Path object or file object.

:param format: Optional format override. If omitted, the format to use is determined from the filename extension. If a file object was used instead of a filename, this parameter should always be used.

:param params: Extra parameters to the image writer.

:returns: None

:exception ValueError: If the output format could not be determined

from the file name. Use the format option to solve this.

:exception IOError: If the file could not be written. The file may have been created, and may contain partial data.

In [18]:

```
# The save method has a couple of parameters which are interesting. The first, called fp, is the filename  
# we want to save the object too. The second, format, is interesting, it allows us to change the type of  
# the image, but the docs tell us that this should be done automatically by looking at the file extension  
# as well. Lets give it a try -- this file was originally a GifImageFile, but I bet if we save it with a  
# .png format and read it in again we'll get a different kind of file  
image.save("msi_recruitment.png")  
image=Image.open("msi_recruitment.png")  
import inspect  
inspect.getmro(type(image))
```

Out[18]:

```
(PIL.PngImagePlugin.PngImageFile,  
 PIL.ImageFile.ImageFile,  
 PIL.Image.Image,  
 object)
```

Using ImageFilter

In [19]:

```
# Indeed, this created a new file, which we could view by going to the Jupyter notebook file list by clicking  
# on the logo at the top of the browser, and we can see this new object is actually a PngImageFile object  
# For the purposes of this class the difference in image formats isn't so important, but it's nice that you can  
# explore how a library works using the functions of help(), dir() and getmro().  
#  
# The PILLOW library also has some nice image filters to add some effects. It does this through the filter()  
# function. The filter() function takes a Filter object, and those are all stored in the ImageFilter object.  
# Lets take a look.  
from PIL import ImageFilter  
help(ImageFilter)
```

Help on module PIL.ImageFilter in PIL:

NAME

PIL.ImageFilter

DESCRIPTION

```
# The Python Imaging Library.
# $Id$
#
# standard filters
#
# History:
# 1995-11-27 fl    Created
# 2002-06-08 fl    Added rank and mode filters
# 2003-09-15 fl    Fixed rank calculation in rank filter; added e
xpend call
#
# Copyright (c) 1997-2003 by Secret Labs AB.
# Copyright (c) 1995-2002 by Fredrik Lundh.
#
# See the README file for information on usage and redistributio
n.
#
```

CLASSES

```
builtins.object
  Filter
    ModeFilter
    MultibandFilter
    BoxBlur
    BuiltinFilter
      BLUR
      CONTOUR
      DETAIL
      EDGE_ENHANCE
      EDGE_ENHANCE_MORE
      EMBOSS
      FIND_EDGES
      Kernel
      SHARPEN
      SMOOTH
      SMOOTH_MORE
    Color3DLUT
    GaussianBlur
    UnsharpMask
    RankFilter
      MaxFilter
      MedianFilter
      MinFilter
```

```
class BLUR(BuiltinFilter)
|   Method resolution order:
|   BLUR
|   BuiltinFilter
|   MultibandFilter
|   Filter
|   builtins.object
|
|   Data and other attributes defined here:
|
|   filterargs = ((5, 5), 16, 0, (1, 1, 1, 1, 1, 1, 0, 0, 0, 1,
```


1, 0, 0, 0...

```
name = 'Blur'
```

Methods inherited from BuiltinFilter:

```
filter(self, image)
```

Data descriptors inherited from Filter:

```
__dict__
    dictionary for instance variables (if defined)
```

```
__weakref__
    list of weak references to the object (if defined)
```

```
class BoxBlur(MultibandFilter)
```

```
    BoxBlur(radius)
```

Blurs the image by setting each pixel to the average value of the pixels

in a square box extending radius pixels in each direction.

Supports float radius of arbitrary size. Uses an optimized implementation

which runs in linear time relative to the size of the image for any radius value.

:param radius: Size of the box in one direction. Radius 0 does not blur,

returns an identical image. Radius 1 takes 1 pixel

in each direction, i.e. 9 pixels in total.

Method resolution order:

```
BoxBlur
MultibandFilter
Filter
builtins.object
```

Methods defined here:

```
__init__(self, radius)
    Initialize self. See help(type(self)) for accurate signature.
```

```
filter(self, image)
```

Data and other attributes defined here:

```
name = 'BoxBlur'
```

Data descriptors inherited from Filter:

```

|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)
|
class BuiltinFilter(MultibandFilter)
|   Method resolution order:
|       BuiltinFilter
|       MultibandFilter
|       Filter
|       builtins.object
|
|   Methods defined here:
|
|   filter(self, image)
|
|   -----
-----
|   Data descriptors inherited from Filter:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)
|
class CONTOUR(BuiltinFilter)
|   Method resolution order:
|       CONTOUR
|       BuiltinFilter
|       MultibandFilter
|       Filter
|       builtins.object
|
|   Data and other attributes defined here:
|
|   filterargs = ((3, 3), 1, 255, (-1, -1, -1, -1, 8, -1, -1, -
1, -1))
|
|   name = 'Contour'
|
|   -----
-----
|   Methods inherited from BuiltinFilter:
|
|   filter(self, image)
|
|   -----
-----
|   Data descriptors inherited from Filter:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)
|
class Color3DLUT(MultibandFilter)
|   Color3DLUT(size, table, channels=3, target_mode=None, **kwar
gs)

```

```

    Three-dimensional color lookup table.

    Transforms 3-channel pixels using the values of the channels
as coordinates
    in the 3D lookup table and interpolating the nearest element
s.

    This method allows you to apply almost any color transformat
ion
    in constant time by using pre-calculated decimated tables.

    .. versionadded:: 5.2.0

    :param size: Size of the table. One int or tuple of (int, in
t, int).
        Minimal size in any dimension is 2, maximum is
65.
    :param table: Flat lookup table. A list of ``channels * size
**3``
        float elements or a list of ``size**3`` channe
ls-sized
        tuples with floats. Channels are changed firs
t,
        then first dimension, then second, then third.
        Value 0.0 corresponds lowest value of output,
1.0 highest.
    :param channels: Number of channels in the table. Could be 3
or 4.
        Default is 3.
    :param target_mode: A mode for the result image. Should have
not less
        than ``channels`` channels. Default is ``
None``,
        which means that mode wouldn't be change
d.

    Method resolution order:
        Color3DLUT
        MultibandFilter
        Filter
        builtins.object

    Methods defined here:

    __init__(self, size, table, channels=3, target_mode=None, **
kwargs)
        Initialize self. See help(type(self)) for accurate sign
ature.

    __repr__(self)
        Return repr(self).

    filter(self, image)

    transform(self, callback, with_normals=False, channels=None,
target_mode=None)
        Transforms the table values using provided callback and
returns
        a new LUT with altered values.

```

```

values      |      :param callback: A function which takes old lookup table
            |      and returns a new set of values. The nu
mber        |      of arguments which function should take
            |      is
            |      ``self.channels`` or ``3 + self.channel
s``         |      if ``with_normals`` flag is set.
            |      Should return a tuple of ``self.channel
s`` or      |      ``channels`` elements if it is set.
            |      :param with_normals: If true, ``callback`` will be calle
d with      |      coordinates in the color cube as th
e first     |      three arguments. Otherwise, ``callb
ack``       |      will be called only with actual col
or values.  |
            |      :param channels: The number of channels in the resulting
lookup table. |
            |      :param target_mode: Passed to the constructor of the res
ulting      |      lookup table.

```

```

-----
            |      Class methods defined here:
            |
            |      generate(size, callback, channels=3, target_mode=None) from
builtins.type |      Generates new LUT using provided callback.
            |
            |      :param size: Size of the table. Passed to the constructo
r.           |
            |      :param callback: Function with three parameters which co
rrespond     |      three color channels. Will be called ``
size**3``    |      times with values from 0.0 to 1.0 and s
hould return |      a tuple with ``channels`` elements.
            |      :param channels: The number of channels which should ret
urn callback. |
            |      :param target_mode: Passed to the constructor of the res
ulting      |      lookup table.

```

```

-----
            |      Data and other attributes defined here:
            |
            |      name = 'Color 3D LUT'
            |
            |      -----

```

```

-----
            |      Data descriptors inherited from Filter:

```

```

            |      __dict__
            |      dictionary for instance variables (if defined)

```

```

    |   __weakref__
    |       list of weak references to the object (if defined)
    |
class DETAIL(BuiltinFilter)
    |   Method resolution order:
    |       DETAIL
    |       BuiltinFilter
    |       MultibandFilter
    |       Filter
    |       builtins.object
    |
    |   Data and other attributes defined here:
    |
    |   filterargs = ((3, 3), 6, 0, (0, -1, 0, -1, 10, -1, 0, -1,
0))
    |
    |   name = 'Detail'
    |
    |   -----
    |
    |   Methods inherited from BuiltinFilter:
    |
    |   filter(self, image)
    |
    |   -----
    |
    |   Data descriptors inherited from Filter:
    |
    |   __dict__
    |       dictionary for instance variables (if defined)
    |
    |   __weakref__
    |       list of weak references to the object (if defined)
    |
class EDGE_ENHANCE(BuiltinFilter)
    |   Method resolution order:
    |       EDGE_ENHANCE
    |       BuiltinFilter
    |       MultibandFilter
    |       Filter
    |       builtins.object
    |
    |   Data and other attributes defined here:
    |
    |   filterargs = ((3, 3), 2, 0, (-1, -1, -1, -1, 10, -1, -1, -1,
-1))
    |
    |   name = 'Edge-enhance'
    |
    |   -----
    |
    |   Methods inherited from BuiltinFilter:
    |
    |   filter(self, image)
    |
    |   -----
    |
    |   Data descriptors inherited from Filter:
    |
    |   __dict__

```

```

        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)

class EDGE_ENHANCE_MORE(BuiltinFilter)
    Method resolution order:
        EDGE_ENHANCE_MORE
        BuiltinFilter
        MultibandFilter
        Filter
        builtins.object

    Data and other attributes defined here:

    filterargs = ((3, 3), 1, 0, (-1, -1, -1, -1, 9, -1, -1, -1,
-1))

    name = 'Edge-enhance More'

    -----

    Methods inherited from BuiltinFilter:

    filter(self, image)

    -----

    Data descriptors inherited from Filter:

    __dict__
        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)

class EMBOSS(BuiltinFilter)
    Method resolution order:
        EMBOSS
        BuiltinFilter
        MultibandFilter
        Filter
        builtins.object

    Data and other attributes defined here:

    filterargs = ((3, 3), 1, 128, (-1, 0, 0, 0, 1, 0, 0, 0, 0))

    name = 'Emboss'

    -----

    Methods inherited from BuiltinFilter:

    filter(self, image)

    -----

    Data descriptors inherited from Filter:

    __dict__

```

```

        dictionary for instance variables (if defined)
    __weakref__
        list of weak references to the object (if defined)

class FIND_EDGES(BuiltinFilter)
    Method resolution order:
        FIND_EDGES
        BuiltinFilter
        MultibandFilter
        Filter
        builtins.object

    Data and other attributes defined here:

    filterargs = ((3, 3), 1, 0, (-1, -1, -1, -1, 8, -1, -1, -1,
-1))

    name = 'Find Edges'

    -----
    Methods inherited from BuiltinFilter:

    filter(self, image)

    -----
    Data descriptors inherited from Filter:

    __dict__
        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)

class Filter(builtins.object)
    Data descriptors defined here:

    __dict__
        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)

class GaussianBlur(MultibandFilter)
    GaussianBlur(radius=2)

    Gaussian blur filter.

    :param radius: Blur radius.

    Method resolution order:
        GaussianBlur
        MultibandFilter
        Filter
        builtins.object

    Methods defined here:

    __init__(self, radius=2)

```

```

        Initialize self. See help(type(self)) for accurate signature.

        filter(self, image)
        -----
Data and other attributes defined here:

    name = 'GaussianBlur'
    -----
Data descriptors inherited from Filter:

    __dict__
        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)

class Kernel(BuiltinFilter)
    Kernel(size, kernel, scale=None, offset=0)

    Create a convolution kernel. The current version only supports 3x3 and 5x5 integer and floating point kernels.

    In the current version, kernels can only be applied to "L" and "RGB" images.

    :param size: Kernel size, given as (width, height). In the current version, this must be (3,3) or (5,5).
    :param kernel: A sequence containing kernel weights.
    :param scale: Scale factor. If given, the result for each pixel is divided by this value. the default is the sum of the kernel weights.
    :param offset: Offset. If given, this value is added to the result, after it has been divided by the scale factor.

    Method resolution order:
        Kernel
        BuiltinFilter
        MultibandFilter
        Filter
        builtins.object

    Methods defined here:

    __init__(self, size, kernel, scale=None, offset=0)
        Initialize self. See help(type(self)) for accurate signature.
    -----
Data and other attributes defined here:

```



```

|   name = 'Kernel'
|   -----
-----
|   Methods inherited from BuiltinFilter:
|
|   filter(self, image)
|   -----
-----
|   Data descriptors inherited from Filter:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)
|
class MaxFilter(RankFilter)
|   MaxFilter(size=3)
|
|   Create a max filter.  Picks the largest pixel value in a win
dow with the
|   given size.
|
|   :param size: The kernel size, in pixels.
|
|   Method resolution order:
|       MaxFilter
|       RankFilter
|       Filter
|       builtins.object
|
|   Methods defined here:
|
|   __init__(self, size=3)
|       Initialize self.  See help(type(self)) for accurate sign
ature.
|
|   -----
-----
|   Data and other attributes defined here:
|
|   name = 'Max'
|   -----
-----
|   Methods inherited from RankFilter:
|
|   filter(self, image)
|   -----
-----
|   Data descriptors inherited from Filter:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)

```

```

class MedianFilter(RankFilter)
|   MedianFilter(size=3)
|
|   Create a median filter. Picks the median pixel value in a wi
ndow with the
|   given size.
|
|   :param size: The kernel size, in pixels.
|
|   Method resolution order:
|       MedianFilter
|       RankFilter
|       Filter
|       builtins.object
|
|   Methods defined here:
|
|   __init__(self, size=3)
|       Initialize self. See help(type(self)) for accurate sign
ature.
|
|   -----
|
|   Data and other attributes defined here:
|
|   name = 'Median'
|
|   -----
|
|   Methods inherited from RankFilter:
|
|   filter(self, image)
|
|   -----
|
|   Data descriptors inherited from Filter:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)
|
class MinFilter(RankFilter)
|   MinFilter(size=3)
|
|   Create a min filter. Picks the lowest pixel value in a wind
ow with the
|   given size.
|
|   :param size: The kernel size, in pixels.
|
|   Method resolution order:
|       MinFilter
|       RankFilter
|       Filter
|       builtins.object
|
|   Methods defined here:
|
|   __init__(self, size=3)

```

```

|         Initialize self.  See help(type(self)) for accurate sign
ature.
|
|-----
|
| Data and other attributes defined here:
|
| name = 'Min'
|
|-----
|
| Methods inherited from RankFilter:
|
| filter(self, image)
|
|-----
|
| Data descriptors inherited from Filter:
|
| __dict__
|     dictionary for instance variables (if defined)
|
| __weakref__
|     list of weak references to the object (if defined)
|
class ModeFilter(Filter)
|     ModeFilter(size=3)
|
|     Create a mode filter. Picks the most frequent pixel value in
a box with the
|     given size.  Pixel values that occur only once or twice are
ignored; if no
|     pixel value occurs more than twice, the original pixel value
is preserved.
|
| :param size: The kernel size, in pixels.
|
| Method resolution order:
|     ModeFilter
|     Filter
|     builtins.object
|
| Methods defined here:
|
| __init__(self, size=3)
|     Initialize self.  See help(type(self)) for accurate sign
ature.
|
| filter(self, image)
|
|-----
|
| Data and other attributes defined here:
|
| name = 'Mode'
|
|-----
|
| Data descriptors inherited from Filter:
|
| __dict__

```

```

        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)

class MultibandFilter(Filter)
    Method resolution order:
        MultibandFilter
        Filter
        builtins.object

    Data descriptors inherited from Filter:

    __dict__
        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)

class RankFilter(Filter)
    RankFilter(size, rank)

    Create a rank filter.  The rank filter sorts all pixels in
    a window of the given size, and returns the **rank**'th valu
e.

    :param size: The kernel size, in pixels.
    :param rank: What pixel value to pick.  Use 0 for a min filt
er,
                  ``size * size / 2`` for a median filter, ``size
* size - 1``
                  for a max filter, etc.

    Method resolution order:
        RankFilter
        Filter
        builtins.object

    Methods defined here:

    __init__(self, size, rank)
        Initialize self.  See help(type(self)) for accurate sign
ature.

    filter(self, image)

    -----
-----
    Data and other attributes defined here:

    name = 'Rank'

    -----
-----
    Data descriptors inherited from Filter:

    __dict__
        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)

```

```

class SHARPEN(BuiltinFilter)
|   Method resolution order:
|       SHARPEN
|       BuiltinFilter
|       MultibandFilter
|       Filter
|       builtins.object
|
|   Data and other attributes defined here:
|
|   filterargs = ((3, 3), 16, 0, (-2, -2, -2, -2, 32, -2, -2, -
2, -2))
|
|   name = 'Sharpen'
|
|   -----
-----
|   Methods inherited from BuiltinFilter:
|
|   filter(self, image)
|
|   -----
-----
|   Data descriptors inherited from Filter:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)
|
class SMOOTH(BuiltinFilter)
|   Method resolution order:
|       SMOOTH
|       BuiltinFilter
|       MultibandFilter
|       Filter
|       builtins.object
|
|   Data and other attributes defined here:
|
|   filterargs = ((3, 3), 13, 0, (1, 1, 1, 1, 5, 1, 1, 1, 1))
|
|   name = 'Smooth'
|
|   -----
-----
|   Methods inherited from BuiltinFilter:
|
|   filter(self, image)
|
|   -----
-----
|   Data descriptors inherited from Filter:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)

```

```

class SMOOTH_MORE(BuiltinFilter)
|   Method resolution order:
|       SMOOTH_MORE
|       BuiltinFilter
|       MultibandFilter
|       Filter
|       builtins.object
|
|   Data and other attributes defined here:
|
|   filterargs = ((5, 5), 100, 0, (1, 1, 1, 1, 1, 1, 5, 5, 5, 1,
1, 5, 44,...
|
|   name = 'Smooth More'
|
|   -----

```

```

-----
|   Methods inherited from BuiltinFilter:
|
|   filter(self, image)
|
|   -----

```

```

-----
|   Data descriptors inherited from Filter:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)

```

```

class UnsharpMask(MultibandFilter)
|   UnsharpMask(radius=2, percent=150, threshold=3)
|
|   Unsharp mask filter.
|
|   See Wikipedia's entry on `digital unsharp masking`_ for an e
xplanation of
|   the parameters.
|
|   :param radius: Blur Radius
|   :param percent: Unsharp strength, in percent
|   :param threshold: Threshold controls the minimum brightness
change that
|       will be sharpened
|
|   .. _digital unsharp masking: https://en.wikipedia.org/wiki/U
nsharp_masking#Digital_unsharp_masking

```

```

|   Method resolution order:
|       UnsharpMask
|       MultibandFilter
|       Filter
|       builtins.object
|
|   Methods defined here:
|
|   __init__(self, radius=2, percent=150, threshold=3)
|       Initialize self. See help(type(self)) for accurate sign
ature.

```

```
| filter(self, image)
```

```
-----  
| Data and other attributes defined here:
```

```
| name = 'UnsharpMask'
```

```
-----  
| Data descriptors inherited from Filter:
```

```
| __dict__  
|     dictionary for instance variables (if defined)
```

```
| __weakref__  
|     list of weak references to the object (if defined)
```

DATA

```
    division = _Feature((2, 2, 0, 'alpha', 2), (3, 0, 0, 'alpha',  
0), 8192...
```

FILE

```
    /opt/conda/lib/python3.7/site-packages/PIL/ImageFilter.py
```

In [22]:

```
# There are a bunch of different filters here, but lets just try and apply the BLUR filter. Before we do this
# we have to convert the image to RGB mode. This is a bit magical -- images like gifs are limited in how many
# colors can be displayed at once based on the size of the pallet. This is similar to a painters pallet, which
# only has so much room. This is actually a very old image file format. If we convert the image into something
# more sophisticated we can apply these interesting image transforms. Sometimes learning a new library means
# digging a bit deeper into the domain the library is about. We can convert the image using the convert()
# function.
image=image.convert('RGB') # this stands for red, green blue mode
blurred_image=image.filter(PIL.ImageFilter.CONTOUR)
display(blurred_image)
```



Using `crop()` to crop an Image

In [23]:

```
# Ok, let me show you one more function in the lecture, which is crop(). This re  
moves portions of the image  
# except for the bounding box you describe. When you think of images, think of i  
ndividual dots or pixels  
# which make up that image being lined up in a grid. You can actually see the nu  
mber of pixels high the image  
# is and the width of the image  
print("{}x{}".format(image.width, image.height))
```

800x450

In [24]:

```
# This means that the image is 800 pixels wide (the X axis), and 450 pixels high  
(the Y axis). If we take a  
# look at the crop documentation we see that the first parameter to the function  
is a tuple which is the  
# left, upper, right, and lower values of the X/Y coordinates  
help(image.crop)
```

Help on method crop in module PIL.Image:

crop(box=None) method of PIL.Image.Image instance

Returns a rectangular region from this image. The box is a 4-tuple defining the left, upper, right, and lower pixel coordinate. See :ref:`coordinate-system`.

Note: Prior to Pillow 3.4.0, this was a lazy operation.

:param box: The crop rectangle, as a (left, upper, right, lower) tuple.

:rtype: :py:class:`~PIL.Image.Image`

:returns: An :py:class:`~PIL.Image.Image` object.

In [25]:

```
# With PIL images, we define the bounding box using the upper left corner and the  
lower right corner. And  
# we count the number of pixels out from the upper left corner, which is 0,0. This  
might seem odd if you're  
# used to coordinate systems where you start in the lower left -- just remember  
that we define our box in the  
# same way we count out positions in the image.  
#  
# So, if we wanted to get the Michigan logo out of this image, we might start with  
the left at, say 50 pixels,  
# and the top at 0 pixels, then we might walk to the right another 190 pixels, and  
set the lower bound to say  
# 150 pixels  
display(image.crop((50,0,190,150)))
```



Using ImageDraw

In [26]:

```
# Of course crop(), like other functions, only returns a copy of the image, and
# doesn't change the image itself.
# A strategy I like to do is try and draw the bounding box directly on the image,
# when I'm trying to line things
# up. We can draw on images using the ImageDraw object. I'm not going to go into
# this in detail, but here's a
# quick example of how. I might draw the bounding box in this case.
from PIL import ImageDraw
drawing_object=ImageDraw.Draw(image)
drawing_object.rectangle((50,0,190,150), fill = None, outline = 'red')
display(image)
```



Ok, that's been an overview of how to use PIL for single images. But, a lot of work might involve multiple images, and putting images together. In the next lecture we'll tackle that, and set you up for the assignment.

Additional PILLOW functions

Lets take a look at some other functions we might want to use in PILLOW to modify images.

In [28]:

```
# First, lets import all of the library functions we need
import PIL
from PIL import Image
from IPython.display import display
import inspect
# And lets load the image we were working, and we can just convert it to RGB inline
file="readonly/msi_recruitment.gif"
image=Image.open(file).convert('RGB')

display(image)
```



Modifying Brightness

In [29]:

```
# A task that is fairly common in image and picture manipulation is to create contact sheets of images.
# A contact sheet is one image that actually contains several other different images. Lets try and make
# a contact sheet for the Master of Science in Information advertisement image. In particular, lets change
# the brightness of the image in ten different ways, then scale the image down smaller, and put them side
# by side so we can get the sense of which brightness we might want to use.
#
# First up, lets import the ImageEnhance module, which has a nice object called Brightness
from PIL import ImageEnhance
# Checking the online documentation for this function, it takes a value between 0.0 (a completely black
# image) and 1.0 (the original image) to adjust the brightness. All of the classes in the ImageEnhance module
# do this the same way, you create an object, in this case Brightness, then you call the enhance function()
# on that object with an appropriate parameter.
#
# Lets write a little loop to generate ten images of different brightness. First we need the Brightness
# object with our image
enhancer=ImageEnhance.Brightness(image)
images=[]
for i in range(0, 10):
    # We'll divide i by ten to get the decimal value we want, and append it to the images list
    # we actually call the brightness routine by calling the enhance() function. Remember, you can dig into
    # details of this using the help() function, or by consulting web docs
    images.append(enhancer.enhance(i/10))
# We can see the result here is a list of ten PIL.Image.Image objects. Jupyter nicely prints out the value
# of python objects nested in lists
print(images)
```

```
[<PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1B0C18>, <PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1B0DA0>, <PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1B0EF0>, <PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1B0F60>, <PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1C1048>, <PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1C12B0>, <PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1C1320>, <PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1C1390>, <PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1C1400>, <PIL.Image.Image image mode=RGB size=800x450 at 0x7FEC2C1C1470>]
```

In [30]:

```
# Lets take these images now and composite them, one above another, in a contact sheet.  
# There are several different approaches we can use, but I'll simply create a new image which is like  
# the first image, but ten times as high. Lets check out the PIL.Image.new functionality  
help(PIL.Image.new)
```

Help on function new in module PIL.Image:

new(mode, size, color=0)

Creates a new image with the given mode and size.

:param mode: The mode to use for the new image. See:

:ref:`concept-modes`.

:param size: A 2-tuple, containing (width, height) in pixels.

:param color: What color to use for the image. Default is black.

k.

If given, this should be a single integer or floating point value

value

for single-band modes, and a tuple for multi-band modes (one

value

per band). When creating RGB images, you can also use color strings as supported by the ImageColor module. If the color

is

None, the image is not initialised.

:returns: An :py:class:`~PIL.Image.Image` object.

In [31]:

```
# The new function requires that we pass it a mode. We're going to use the mode
# 'RGB' which stands for
# Red, Green, and Blue, and is the mode of our current first image. There are lots
# of different image mode
# formats, and this one is most common.
# For the size we have a tuple, which is the width of the image and the height.
# We'll use the width of our
# current first image, but for the height we'll multiple this by ten. This will
# make a sort of "canvas" for
# our contact sheet. Finally, the color is optional, and we'll just leave it at
# black.
first_image=images[0]
from PIL import Image
contact_sheet=PIL.Image.new(first_image.mode, (first_image.width,10*first_image.
height))

# So now we have a black image that's ten times the size of the other images in
# the contact_sheet
# variable. Now lets just loop through the image list and paste() the results i
# n. The paste() function
# will be called on the contact_sheet object, and takes in a new image to paste,
# as well as an (x,y)
# offset for that image. In our case, the x position is always 0, but the y loca
# tion will change by
# 450 pixels each time we iterate through the loop.
#
# Lets first create a counter variable for the y location. It will start at zero
current_location = 0
for img in images:
    # Lets paste the current image into the contact sheet
    contact_sheet.paste(img, (0, current_location) )
    # And update the current_location counter
    current_location=current_location+450

# This contact sheet has gotten big: 4,500 pixels tall! Lets just resize this sh
# eet for display. We can do
# this using the resize() function. This function just takes a tuple of width an
# d height, and we'll resize
# everything down to the size of just two individual images
contact_sheet = contact_sheet.resize((160,900) )
# Now lets just display that composite image
display(contact_sheet)
```

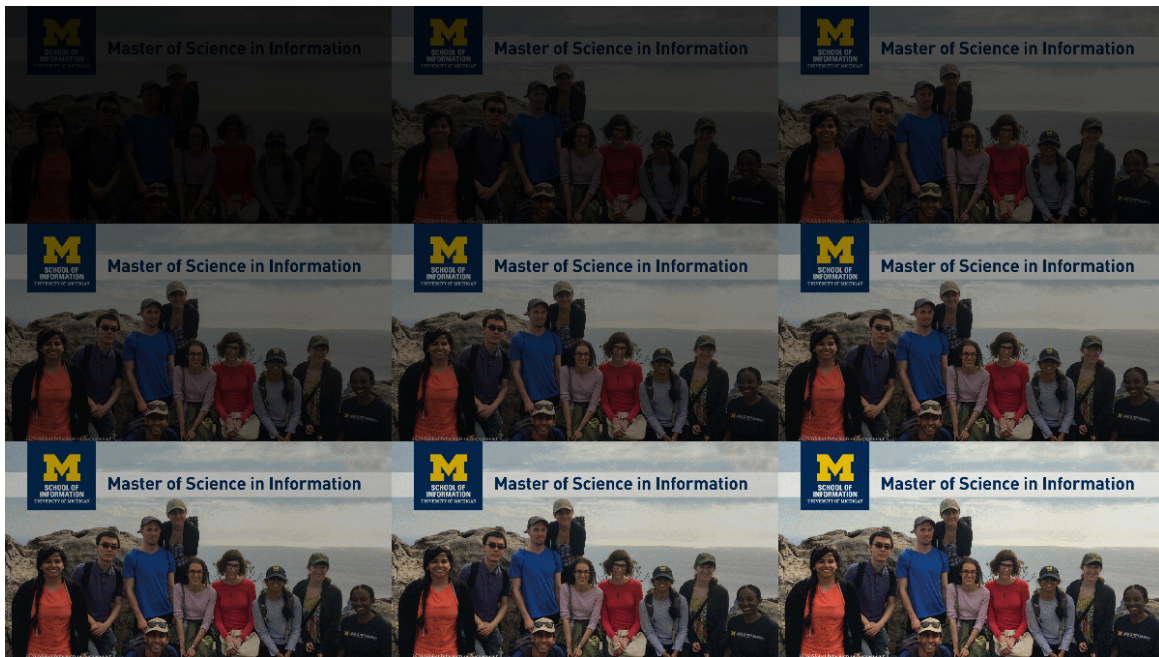


In [32]:

```
# Ok, that's a nice proof of concept. But it's a little tough to see. Lets inste
ad change this to a three
# by three grid of values. First thing we should do is make our canvas, and we'l
l make it 3 times the
# width of our image and 3 times the height of our image - a nine image square
contact_sheet=PIL.Image.new(first_image.mode, (first_image.width*3,first_image.h
eight*3))
# Now we want to iterate over our images and place them into this grid. Remember
that in PIL we manage the
# location of where we refer to as an image in the upper right hand corner, so t
his will be 0,0. Lets use
# one variable for the X dimension, and one for the Y dimension.
x=0
y=0

# Now, lets iterate over our images. Except, we don't want to both with the firs
t one, because it is
# just solid black. Instead we want to just deal with the images after the first
one, and that should
# give us nine in total
for img in images[1:]:
    # Lets paste the current image into the contact sheet
    contact_sheet.paste(img, (x, y) )
    # Now we update our X position. If it is going to be the width of the image,
then we set it to 0
    # and update Y as well to point to the next "line" of the contact sheet.
    if x+first_image.width == contact_sheet.width:
        x=0
        y=y+first_image.height
    else:
        x=x+first_image.width

# Now lets resize the contact sheet. We'll just make it half the size by dividin
g it by two. And, because
# the resize function needs to take round numbers, we need to convert our divisi
ons from floating point
# numbers into integers using the int() function.
contact_sheet = contact_sheet.resize((int(contact_sheet.width/2),int(contact_she
et.height/2) ))
# Now lets display that composite image
display(contact_sheet)
```



Well, that's been a tour of our first external API, the Python Imaging Library, or pillow module. In this series of lectures you've learned how to read and write images, manipulate them with pillow, and explore the functionality of third party APIs using features of Python like `dir()`, `help()`, and `getmro()`. You've also been introduced to the console, and how python stores these libraries on the computer. While for this course all of the libraries are included for you in the Coursera system, and you won't need to install your own, it's good to get a the idea of how this work in case you wanted to set this up on your own.

Finally, while you can explore PILLOW from within python, most good modules also put their documentation up online, and you can read more about PILLOW here: <https://pillow.readthedocs.io/en/latest/> (<https://pillow.readthedocs.io/en/latest/>)