

03 Computer Vision With OpenCV

The next library we're going to look at is called Kraken, which was developed by Université PSL in Paris. It's actually based on a slightly older code base, **OCROPUS**. You can see how the flexible open-source licenses allow new ideas to grow by building upon older ideas. And, in this case, I fully support the idea that the Kraken - a mythical massive sea creature - is the natural progression of an octopus!

03-01: Release the Kraken!

Purpose of Kraken*

What we are going to use Kraken for is to **detect lines of text as bounding boxes** in a given image. The biggest limitation of tesseract is the *lack of a layout engine inside of it*. Tesseract expects to be using fairly clean text, and gets confused if we don't crop out other artifacts. It's not bad, but Kraken can help us out by segmenting pages. Let's take a look.

*Please note that Kraken is only supported on Linux and Mac OS X, it is not supported on Windows. Documentation and Installation Notes can be found at: <https://pypi.org/project/kraken/>

Importing Kraken Module

In [1]:

```
# First, we'll take a look at the kraken module itself
import kraken
help(kraken)
```

Help on package kraken:

NAME

kraken - entry point for kraken functionality

PACKAGE CONTENTS

binarization
ketos
kraken
lib (package)
linegen
pageseg
repo
rpred
serialization
transcribe

DATA

```
absolute_import = _Feature((2, 5, 0, 'alpha', 1), (3, 0, 0, 'alp
ha', 0)...
division = _Feature((2, 2, 0, 'alpha', 2), (3, 0, 0, 'alpha',
0), 8192)...
print_function = _Feature((2, 6, 0, 'alpha', 2), (3, 0, 0, 'alph
a', 0)...
```

FILE

/opt/conda/lib/python3.7/site-packages/kraken/__init__.py

Note the package contents in Kraken.

Importing `pageseg`

There isn't much of a discussion here, but there are a number of sub-modules that look interesting. I spend a bit of time on their website, and I think the `pageseg` module, which handles all of the page segmentation, is the one we want to use.

In [2]:

```
from kraken import pageseg  
help(pageseg)
```

Help on module kraken.pageseg in kraken:

NAME

kraken.pageseg

DESCRIPTION

kraken.pageseg

~~~~~

Layout analysis and script detection methods.

FUNCTIONS

detect\_scripts(im, bounds, model='/opt/conda/lib/python3.7/site-packages/kraken/script.mlmodel', valid\_scripts=None)

Detects scripts in a segmented page.

Classifies lines returned by the page segmenter into runs of scripts/writing systems.

Args:

im (PIL.Image): A bi-level page of mode '1' or 'L'

bounds (dict): A dictionary containing a 'boxes' entry with a list of coordinates (x0, y0, x1, y1) of a text line in the image

and an entry 'text\_direction' containing 'horizontal-lr/rl/vertical-lr/rl'.

model (str): Location of the script classification model or None for default.

valid\_scripts (list): List of valid scripts.

Returns:

{'script\_detection': True, 'text\_direction': '\$dir', 'boxes':

[[('script', (x1, y1, x2, y2)),...]]}: A dictionary containing the text direction and a list of lists of reading order sorted bounding boxes under the key 'boxes' with each list containing the script segmentation of a single line. Script is a ISO15924 4 character identifier.

Raises:

KrakenInvalidModelError if no clstm module is available.

segment(im, text\_direction='horizontal-lr', scale=None, maxcolseps=2, black\_colseps=False, no\_hlines=True, pad=0, mask=None)

Segments a page into text lines.

Segments a page into text lines and returns the absolute coordinates of each line in reading order.

Args:

im (PIL.Image): A bi-level page of mode '1' or 'L'

text\_direction (str): Principal direction of the text (horizontal-lr/rl/vertical-lr/rl)

scale (float): Scale of the image

maxcolseps (int): Maximum number of whitespace column se

```

parators
    black_colseps (bool): Whether column separators are assumed to be
                           vertical black lines or not
    no_hlines (bool): Switch for horizontal line removal
    pad (int or tuple): Padding to add to line bounding boxes. If int the
                        same padding is used both left and right. If a
                        2-tuple, uses (padding_left, padding_right).
    mask (PIL.Image): A bi-level mask image of the same size as `im` where
                      0-valued regions are ignored for segmentation
                      purposes. Disables column detection.

Returns:
    {'text_direction': '$dir', 'boxes': [(x1, y1, x2, y2), ...]}: A
    dictionary containing the text direction and a list of reading order
    sorted bounding boxes under the key 'boxes'.

Raises:
    KrakenInputException if the input image is not binarized or the text
    direction is invalid.

DATA
    __all__ = ['segment', 'detect_scripts']

FILE
    /opt/conda/lib/python3.7/site-packages/kraken/pageseg.py

```

So it looks like there are a few different functions we can call, and the segment function looks particularly appropriate.

The `kraken.pageseg` library is on the documentation front -- I can see immediately that we are working with *PIL.Image files*, and the author has even indicated that we need to pass in **either a binarized (e.g. '1') or grayscale (e.g. 'L') image**.

We can also see that the **return value is a dictionary object with two keys**, "text\_direction" which will return to us...

- a string of the direction of the text,
- and "boxes" which appears to be a list of tuples,

where each tuple is a box in the original image.

## Importing A MultiColumn Image With Text

This is an excerpt taken from an on-campus draft paper from the University of Michigan.

In [3]:

```
from PIL import Image
im=Image.open("readonly/two_col.png")
# Lets display the image inline
display(im)
# Lets now convert it to black and white and segment it up into lines with krake
n
bounding_boxes=pageseg.segment(im.convert('1'))['boxes']
# And lets print those lines to the screen
print(bounding_boxes)
```

---

**CALEB CHADWELL***Daily Staff Reporter*

---

In a statement Tuesday, the Department of Public Safety and Security wrote that DPSS was not aware until Saturday afternoon of an assault on a University lecturer last week, referred in testimony before Ann Arbor City Council Monday.

Khita Whyatt, lecturer of dance in the University of Michigan's School of Music, Theatre & Dance, said in an interview after her testimony on the incident that she did not immediately call the police because she was so shocked, but her department chair contacted the DPSS. Two days after the incident, Whyatt said she was interviewed by two DPSS officers. During her testimony to Council, she called on the University to release a crime report about how she was knocked

down and intimidated by unknown assailants. The event follows similar incidents where crime alerts had not been released. The University has released two crime alerts of hate crimes on campus over the past two weeks.

Whyatt wrote in an email sent Tuesday afternoon to recipients including University President Mark Schlissel as well as The Michigan Daily that she waited until Saturday morning to report the assault to police because she was disoriented and did not know where to reach out.

"I did wait until Saturday morning to get in touch to report the incident," Whyatt wrote. "I was in shock and still processing what to do prior to reaching out ... it was also obvious that there was no way that these boys were going to be caught. Not being a student, I did not know who to report to. That must seem obvious by the fact that I

```
[[100, 50, 449, 74], [131, 88, 414, 120], [59, 196, 522, 229], [18, 239, 522, 272], [19, 283, 522, 316], [19, 327, 525, 360], [19, 371, 523, 404], [18, 414, 524, 447], [17, 458, 522, 491], [19, 502, 141, 535], [58, 546, 521, 579], [18, 589, 522, 622], [19, 633, 521, 665], [563, 21, 1066, 54], [564, 64, 1066, 91], [563, 108, 1066, 135], [564, 152, 1065, 179], [563, 196, 1065, 229], [563, 239, 1066, 272], [562, 283, 909, 316], [600, 327, 1066, 360], [562, 371, 1066, 404], [562, 414, 1066, 447], [563, 458, 1065, 485], [563, 502, 1065, 535], [562, 546, 1066, 579], [562, 589, 1064, 622], [562, 633, 1066, 660], [18, 677, 833, 704], [18, 721, 1066, 754], [18, 764, 1065, 797], [17, 808, 1065, 841], [18, 852, 1067, 885], [18, 895, 1065, 928], [17, 939, 1065, 972], [17, 983, 1067, 1016], [18, 1027, 1065, 1060], [18, 1070, 1065, 1103], [18, 1114, 1065, 1147]]
```

## Printing Bounding Boxes To Show Segregation

Let's write a little routine to try and see the effects a bit more clearly. Add a little documentation using triple quoted strings `''' '''`. Basically, this series of code will colour the boxes red.

In [4]:

```
def show_boxes(img):  
    '''Modifies the passed image to show a series of bounding boxes on an image  
    as run by kraken  
  
    :param img: A PIL.Image object  
    :return img: The modified PIL.Image object  
    '''  
  
    # Lets bring in our ImageDraw object  
    from PIL import ImageDraw  
    # And grab a drawing object to annotate that image  
    drawing_object=ImageDraw.Draw(img)  
    # We can create a set of boxes using pageseg.segment  
    bounding_boxes=pageseg.segment(img.convert('1'))['boxes']  
    # Now lets go through the list of bounding boxes  
    for box in bounding_boxes:  
        # An just draw a nice rectangle  
        drawing_object.rectangle(box, fill = None, outline ='red')  
    # And to make it easy, lets return the image object  
    return img  
  
# To test this, lets use display  
display(show_boxes(Image.open("readonly/two_col.png")))
```



---

**CALEB CHADWELL***Daily Staff Reporter*

---

In a statement Tuesday, the Department of Public Safety and Security wrote that DPSS was not aware until Saturday afternoon of an assault on a University lecturer last week, referred in testimony before Ann Arbor City Council Monday.

Khita Whyatt, lecturer of dance in the University of Michigan's School of Music, Theatre & Dance, said in an interview after her testimony on the incident that she did not immediately call the police because she was so shocked, but her department chair contacted the DPSS. Two days after the incident, Whyatt said she was interviewed by two DPSS officers. During her testimony to Council, she called on the University to release a crime report about how she was knocked

down and intimidated by unknown assailants. The event follows similar incidents where crime alerts had not been released. The University has released two crime alerts of hate crimes on campus over the past two weeks.

Whyatt wrote in an email sent Tuesday afternoon to recipients including University President Mark Schlissel as well as The Michigan Daily that she waited until Saturday morning to report the assault to police because she was disoriented and did not know where to reach out.

"I did wait until Saturday morning to get in touch to report the incident," Whyatt wrote. "I was in shock and still processing what to do prior to reaching out ... it was also obvious that there was no way that these boys were going to be caught. Not being a student, I did not know who to report to. That must seem obvious by the fact that I

## Analysis of Kraken's Performance

Kraken isn't completely sure what to do with this two column format. In some cases, kraken has identified a line in just a single column, while in other cases kraken has spanned the line marker all the way across the page.

## Using the `pageseg()` function's `black_colseps` parameter

If the `black_colseps()` is set to `True`, Kraken will assume that columns will be separated by black lines. This isn't the case here, but we have all the tools to go through and actually change the source image to have a black separator between columns.

First, update `show_boxes()` function by adding `black_colseps = True`

In [5]:

```
# The first step is that I want to update the show_boxes() function. I'm just going to do a quick
# copy and paste from the above but add in the black_colseps=True parameter
def show_boxes(img):
    '''Modifies the passed image to show a series of bounding boxes on an image as run by kraken

    :param img: A PIL.Image object
    :return img: The modified PIL.Image object
    '''

    # Lets bring in our ImageDraw object
    from PIL import ImageDraw
    # And grab a drawing object to annotate that image
    drawing_object=ImageDraw.Draw(img)
    # We can create a set of boxes using pageseg.segment
    bounding_boxes=pageseg.segment(img.convert('1'), black_colseps=True) ['boxes'
]
    # Now lets go through the list of bounding boxes
    for box in bounding_boxes:
        # An just draw a nice rectangle
        drawing_object.rectangle(box, fill = None, outline ='red')
    # And to make it easy, lets return the image object
    return img
```

## Second, create an algorithm to detect a white column separator

In experimenting a bit I decided that I only wanted to add the separator if the space of was at least 25 pixels wide, which is roughly the width of a character, and six lines high. In determining these 2 properties, we need to determine:

- the width: which can be found by just assigning a variable to it at about 25 pixels.
- the height: which is harder since it depends on the height of the text. So write a routine to calculate the average height of a line.

In [6]:

```
def calculate_line_height(img):
    '''Calculates the average height of a line from a given image
    :param img: A PIL.Image object
    :return: The average line height in pixels
    '''
    # Lets get a list of bounding boxes for this image
    bounding_boxes=pageseg.segment(img.convert('1'))['boxes']
    # Each box is a tuple of (top, left, bottom, right) so the height is just to
    p - bottom
    # So lets just calculate this over the set of all boxes
    height_accumulator=0
    for box in bounding_boxes:
        height_accumulator=height_accumulator+box[3]-box[1]
        # this is a bit tricky, remember that we start counting at the upper left
        corner in PIL!
    # now lets just return the average height
    # lets change it to the nearest full pixel by making it an integer
    return int(height_accumulator/len(bounding_boxes))

# And lets test this with the image with have been using
line_height=calculate_line_height(Image.open("readonly/two_col.png"))
print(line_height)
```

31

## Creating a Gap Box

Now after finding the line height, we want to scan through the image, look at each pixel to determine if there is a **block of whitespace**. How big of a block should we look for?

An appropriate box would be about 1 char width wide, and 6 lune heights tall, for instance.

In [7]:

```
# Lets create a new box called gap box that represents this area
gap_box=(0,0,char_width,line_height*6)
gap_box
```

Out[7]:

(0, 0, 25, 186)

## Creating Function gap\_check()

It seems we will want to have a function which, given a pixel in an image, can check to see if that pixel has whitespace to the right and below it. **Essentially, we want to test to see if the pixel is the upper left corner of something that looks like the gap\_box.** If so, then we should insert a line to "break up" this box before sending to kraken

In [8]:

```
# Lets call this new function gap_check
def gap_check(img, location):
    '''Checks the img in a given (x,y) location to see if it fits the description
    of a gap_box
    :param img: A PIL.Image file
    :param location: A tuple (x,y) which is a pixel location in that image
    :return: True if that fits the definition of a gap_box, otherwise False
    '''
    # Recall that we can get a pixel using the img.getpixel() function. It returns this value
    # as a tuple of integers, one for each color channel. Our tools all work with binarized
    # images (black and white), so we should just get one value. If the value is 0 it's a black
    # pixel, if it's white then the value should be 255
    #
    # We're going to assume that the image is in the correct mode already, e.g. it has been
    # binarized. The algorithm to check our bounding box is fairly easy: we have a single location
    # which is our start and then we want to check all the pixels to the right of that location
    # up to gap_box[2]
    for x in range(location[0], location[0]+gap_box[2]):
        # the height is similar, so lets iterate a y variable to gap_box[3]
        for y in range(location[1], location[1]+gap_box[3]):
            # we want to check if the pixel is white, but only if we are still within the image
            if x < img.width and y < img.height:
                # if the pixel is white we don't do anything, if it's black, we just want to
                # finish and return False
                if img.getpixel((x,y)) != 255:
                    return False
            # If we have managed to walk all through the gap_box without finding any non-white pixels
            # then we can return true -- this is a gap!
    return True
```

## Drawing the column separator

Alright, we have a function to check for a gap, called `gap_check`. What should we do once we find a gap? For this, let's just draw a line in the middle of it.

This is important as this is where your pillow library comes in once again. Take a close look at the documentation of the function `draw_sep()` by the instructor.

In [9]:

```
# Lets create a new function
def draw_sep(img,location):
    '''Draws a line in img in the middle of the gap discovered at location. Note
    that
    this doesn't draw the line in location, but draws it at the middle of a gap_
    box
    starting at location.
    :param img: A PIL.Image file
    :param location: A tuple(x,y) which is a pixel location in the image
    '''

    # First lets bring in all of our drawing code
    from PIL import ImageDraw
    drawing_object=ImageDraw.Draw(img)
    # next, lets decide what the middle means in terms of coordinates in the ima
    ge
    x1=location[0]+int(gap_box[2]/2)
    # and our x2 is just the same thing, since this is a one pixel vertical line
    x2=x1
    # our starting y coordinate is just the y coordinate which was passed in, th
    e top of the box
    y1=location[1]
    # but we want our final y coordinate to be the bottom of the box
    y2=y1+gap_box[3]
    drawing_object.rectangle((x1,y1,x2,y2), fill = 'black', outline ='black')
    # and we don't have anything we need to return from this, because we modifie
    d the image
```

### Reprocessing the image to include black lines between columns

Now, lets try it all out. This is pretty easy, we can just iterate through each pixel in the image, check if there is a gap, then insert a line if there is.

**Unfortunately, this has a time complexity of  $O(mn)$ ,**

for an image with  $m$  rows and  $n$  columns.

In [10]:

```
def process_image(img):  
    '''Takes in an image of text and adds black vertical bars to break up column  
s  
:param img: A PIL.Image file  
:return: A modified PIL.Image file  
'''  
  
    # we'll start with a familiar iteration process  
    for x in range(img.width):  
        for y in range(img.height):  
            # check if there is a gap at this point  
            if (gap_check(img, (x,y))):  
                # then update image to one which has a separator drawn on it  
                draw_sep(img, (x,y))  
        # and for good measure we'll return the image we modified  
    return img  
  
# Lets read in our test image and convert it through binarization  
i=Image.open("readonly/two_col.png").convert("L")  
i=process_image(i)  
display(i)  
  
# if you realised, this is  $O(n^2)$  compile time.
```

---

**CALEB CHADWELL***Daily Staff Reporter*

---

In a statement Tuesday, the Department of Public Safety and Security wrote that DPSS was not aware until Saturday afternoon of an assault on a University lecturer last week, referred in testimony before Ann Arbor City Council Monday.

Khita Whyatt, lecturer of dance in the University of Michigan's School of Music, Theatre & Dance, said in an interview after her testimony on the incident that she did not immediately call the police because she was so shocked, but her department chair contacted the DPSS. Two days after the incident, Whyatt said she was interviewed by two DPSS officers. During her testimony to Council, she called on the University to release a crime report about how she was knocked

down and intimidated by unknown assailants. The event follows similar incidents where crime alerts had not been released. The University has released two crime alerts of hate crimes on campus over the past two weeks.

Whyatt wrote in an email sent Tuesday afternoon to recipients including University President Mark Schlissel as well as The Michigan Daily that she waited until Saturday morning to report the assault to police because she was disoriented and did not know where to reach out.

"I did wait until Saturday morning to get in touch to report the incident," Whyatt wrote. "I was in shock and still processing what to do prior to reaching out ... it was also obvious that there was no way that these boys were going to be caught. Not being a student, I did not know who to report to. That must seem obvious by the fact that I

**Pass the Image Back Into Kraken**

In [11]:

```
display(show_boxes(i))
```

---

**CALEB CHADWELL**

*Daily Staff Reporter*

---

In a statement Tuesday, the Department of Public Safety and Security wrote that DPSS was not aware until Saturday afternoon of an assault on a University lecturer last week, referred in testimony before Ann Arbor City Council Monday.

Khita Whyatt, lecturer of dance in the University of Michigan's School of Music, Theatre & Dance, said in an interview after her testimony on the incident that she did not immediately call the police because she was so shocked, but her department chair contacted the DPSS. Two days after the incident, Whyatt said she was interviewed by two DPSS officers. During her testimony to Council, she called on the University to release a crime report about how she was knocked

down and intimidated by unknown assailants. The event follows similar incidents where crime alerts had not been released. The University has released two crime alerts of hate crimes on campus over the past two weeks.

Whyatt wrote in an email sent Tuesday afternoon to recipients including University President Mark Schlissel as well as The Michigan Daily that she waited until Saturday morning to report the assault to police because she was disoriented and did not know where to reach out.

"I did wait until Saturday morning to get in touch to report the incident," Whyatt wrote. "I was in shock and still processing what to do prior to reaching out ... it was also obvious that there was no way that these boys were going to be caught. Not being a student, I did not know who to report to. That must seem obvious by the fact that I



Looks like that is pretty accurate, and fixes the problem we faced. Feel free to experiment with different settings for the gap heights and width and share in the forums. You'll notice though method we created is really quite slow, which is a bit of a problem if we wanted to use this on larger text. But I wanted to show you how you can mix your own logic and work with libraries you're using. Just because Kraken didn't work perfectly, doesn't mean we can't build something more specific to our use case on top of it.

I want to end this lecture with a pause and to ask you to reflect on the code we've written here. We started this course with some pretty simple use of libraries, but now we're digging in deeper and solving problems ourselves with the help of these libraries. Before we go on to our last library, how well prepared do you think you are to take your python skills out into the wild?

## 03-02: Comparing Image Data Structures

OpenCV supports reading of images in most file formats, such as JPEG, PNG, and TIFF. Most image and video analysis requires converting images into grayscale first. This simplifies the image and reduces noise allowing for improved analysis.

### Converting an Image To Grayscale using `cv.cvtColor()`

Let's write some code that reads an image of a person, Floyd Mayweather, and converts it into grayscale.

### OpenCVs Documentation

Now, before we get to the result, let's talk about docs. Just like tesseract, opencv is an external package written in C++, and the docs for python are really poor. This is unfortunately quite common when python is being used as a wrapper. Thankfully, the web docs for opencv are actually pretty good, so hit the website [docs.opencv.org](https://docs.opencv.org) when you want to learn more about a particular function. In this case `cvtColor` converts from one color space to another, and we are converting our image to grayscale.

So now, we can change an image to grayscale in 3 ways:

- Using PIL's changing of color conventions
- Using PIL's Binarisation techniques
- Using OpenCV's `_cv.cvtColor()`

In [12]:

```
# First we will import the open cv package cv2
import cv2 as cv
# We'll load the floyd.jpg image
img = cv.imread('readonly/floyd.jpg')
# And we'll convert it to grayscale using the cvtColor image
gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)

# Lets inspect this object that has been returned.
import inspect
inspect.getmro(type(gray))
```

Out[12]:

```
(numpy.ndarray, object)
```

## Introducing the Numpy Library

We see that it is of type ndarray, which is a fundamental list type coming from the numerical python project. That's a bit surprising - up until this point we have been used to working with PIL.Image objects. **OpenCV, however, wants to represent an image as a two dimensional sequence of bytes**, and the ndarray, which stands for *n dimensional array*, is the ideal way to do this. Lets look at the array contents.

In [13]:

```
gray
```

Out[13]:

```
array([[ 40,  39,  39, ...,  77,  76,  75],
       [ 43,  42,  42, ...,  76,  75,  75],
       [ 39,  39,  39, ...,  76,  75,  74],
       ...,
       [ 21,  22,  24, ..., 219, 223, 209],
       [ 18,  20,  22, ..., 196, 206, 196],
       [ 16,  18,  20, ..., 168, 182, 176]], dtype=uint8)
```

The array is shown here as a list of lists, where the inner lists are filled with integers. The `dtype=uint8` definition indicates that each of the items in an array is an 8 bit unsigned integer, which is very common for black and white images.


So this is a pixel by pixel definition of the image.

## Converting Image into PIL Object

PIL can take an array of data with a given color format and convert this into a PIL object. This is perfect for our situation, as the PIL color mode, "L" is just an array of luminance values in unsigned integers

In [14]:

```
from PIL import Image  
  
image = Image.fromarray(gray, "L")  
display(image)
```



## Multidimensional Images: Introduction to Numpy & List of Lists

Lets talk a bit more about images for a moment. Numpy arrays are multidimensional. For instance, we can define an array in a single dimension:

In [ ]:

```
import numpy as np  
single_dim = np.array([25, 50 , 25, 10, 10])
```

In an image, this is analagous to a **single row of 5 pixels each in grayscale**. But actually, all **imaging libraries tend to expect at least two dimensions**, a width and a height, and to show a matrix. So if we put the single\_dim inside of another array, this would be a two dimensional array with element in the height direction, and five in the width direction

In [15]:

```
double_dim = np.array([single_dim])  
  
double_dim
```

Out[15]:

```
array([[25, 50, 25, 10, 10]])
```

This should look pretty familiar, **it's a lot like a list of lists!** Lets see what this new two dimensional array looks like if we display it

In [16]:

```
display(Image.fromarray(double_dim, "L"))
```

-

Pretty unexciting - it's just a little line. Five pixels in a row to be exact, **of different levels of black**. We can check this by looking at the function `np.shape`. The shape attribute **returns a tuple** that shows the height of the image, by the width of the image.

In [17]:

```
# Pretty unexciting - it's just a little line. Five pixels in a row to be exact,  
of different  
# levels of black. The numpy library has a nice attribute called shape that allo  
ws us to see how  
# many dimensions big an array is. The shape attribute returns a tuple that show  
s the height of  
# the image, by the width of the image  
double_dim.shape
```

Out[17]:

```
(1, 5)
```

In [18]:

```
# Lets take a look at the shape of our initial image which we loaded into the im  
g variable  
img.shape
```

Out[18]:

```
(416, 416, 3)
```

In [19]:

```
# This image has three dimensions! That's because it has a width, a height, and
# what's called
# a color depth. In this case, the color is represented as an array of three val
# ues. Lets take a
# look at the color of the first pixel
first_pixel=img[0][0]
first_pixel
```

Out[19]:

```
array([33, 35, 53], dtype=uint8)
```

Here we see that the **color value is provided in full RGB** using an unsigned integer. This means that **each color can have one of 256 values**, and the total number of unique colors that can be represented by this data is 256 256 256 which is roughly 16 million colors. We call this 24 bit color, which is 8+8+8.

## Changing Image into Grayscale via `reshape()`

One of the most common things to do with an ndarray is to reshape it -- to change the number of rows and columns that are represented so that we can do different kinds of operations. Here is our original two dimensional image

In addition, element-wise operations on the computer are done simultaneously, which means such operations are on average  $O(1)$  time.

In [20]:

```
print("Original image")
print(gray)
# If we wanted to represent that as a one dimensional image, we just call reshap
e
print("New image")
# And reshape takes the image as the first parameter, and a new shape as the sec
ond
image1d=np.reshape(gray,(1,gray.shape[0]*gray.shape[1]))
print(image1d)
```

Original image

```
[[ 40  39  39 ...  77  76  75]
 [ 43  42  42 ...  76  75  75]
 [ 39  39  39 ...  76  75  74]
 ...
 [ 21  22  24 ... 219 223 209]
 [ 18  20  22 ... 196 206 196]
 [ 16  18  20 ... 168 182 176]]
```

New image

```
[[ 40  39  39 ... 168 182 176]]
```

So, why are we talking about these nested arrays of bytes, we were supposed to be talking about OpenCV as a library. Well, I wanted to show you that often libraries working on the same kind of principles, in this case images stored as arrays of bytes, are not representing data in the same way in their APIs. But, by exploring a bit you can learn how the internal representation of data is stored, and build routines to convert between formats.

## For instance, remember in the last lecture when we wanted to look for gaps in an image so

that we could draw lines to feed into kraken? Well, we use PIL to do this, using `getpixel()`

## \ to look at individual pixels and see what the luminosity was, then `ImageDraw.rectangle` to

actually fill in a black bar separator. This was a nice high level API, and let us write routines to do the work we wanted without having to understand too much about how the images were being stored. But it was computationally very slow.

### Creating GrayScale using Matrix Operation

Now, remember how slicing on a list works, if you have a list of number such as `a=[0,1,2,3,4,5]` then `a[2:4]` will return the sublist of numbers at position 2 through 4 inclusive - don't forget that lists start indexing at 0!

In [ ]:

```
# Instead, we could write the code to do this using matrix features within numpy. Lets take  
# a look.  
import cv2 as cv  
# We'll load the 2 column image  
img = cv.imread('readonly/two_col.png')  
# And we'll convert it to grayscale using the cvtColor image  
gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
```

In [21]:

```
# Now, remember how slicing on a list works, if you have a list of number such as  
# a=[0,1,2,3,4,5] then a[2:4] will return the sublist of numbers at position 2 t  
hrough 4  
# inclusive - don't forget that lists start indexing at 0!  
# If we have a two dimensional array, we can slice out a smaller piece of that u  
sing the  
# format a[2:4,1:3]. You can think of this as first slicing along the rows dimen  
sion, then  
# in the columns dimension. So in this example, that would be a matrix of rows  
2, and 3,  
# and columns 1, and 2. Here's a look at our image.  
gray[2:4,1:3]
```

Out[21]:

```
array([[39, 39],  
       [37, 37]], dtype=uint8)
```

In [22]:

```
# So we see that it is all white. We can use this as a "window" and move it arou  
nd our  
# our big image.  
#  
# Finally, the ndarray library has lots of matrix functions which are generally  
very fast  
# to run. One that we want to consider in this case is count_nonzero(), which ju  
st returns  
# the number of entries in the matrix which are not zero.  
np.count_nonzero(gray[2:4,1:3])
```

Out[22]:

4

In [23]:

```
# Ok, the last benefit of going to this low level approach to images is that we
    can change
# pixels very fast as well. Previously we were drawing rectangles and setting a
    fill and line
# width. This is nice if you want to do something like change the color of the f
    ill from the
# line, or draw complex shapes. But we really just want a line here. That's real
    ly easy to
# do - we just want to change a number of luminosity values from 255 to 0.
#
# As an example, lets create a big white matrix
white_matrix=np.full((12,12),255,dtype=np.uint8)
display(Image.fromarray(white_matrix,"L"))
white_matrix
```

Out[23]:

[illegible]



```
# looks pretty boring, it's just a giant white square we can't see. But if we wa
nt, we can
# easily color a column to be black
white_matrix[:,6]=np.full((1,12),0,dtype=np.uint8)
display(Image.fromarray(white_matrix,"L"))
white_matrix
```

Out[24]:

```
5]],
dtype=uint8)
```

And that's exactly what we wanted to do. So, why do it this way, when it seems so much more low level? **Really, the answer is speed.** This paradigm of using matrices to store and manipulate bytes of data for images is much closer to how low level API and hardware developers think about storing files and bytes in memory.

OpenCV comes with trained models for detecting faces, eyes, and smiles which we'll be using. You can train models for detecting other things - like hot dogs or flutes - and if you're interested in that I'd recommend you check out the Open CV docs on how to train a cascade classifier:

After these classifiers are loaded, we want to try and detect a face, convert it to grayscale.

After these classifiers are loaded, we want to try and detect a face, convert it to grayscale.

In [25]:

```
# First step is to load opencv and the XML-based classifiers
import cv2 as cv
face_cascade = cv.CascadeClassifier('readonly/haarcascade_frontalface_default.xml')
eye_cascade = cv.CascadeClassifier('readonly/haarcascade_eye.xml')
# Ok, with the classifiers loaded, we now want to try and detect a face. Lets put
# it in the
# picture we played with last time
img = cv.imread('readonly/floyd.jpg')
# And we'll convert it to grayscale using the cvtColor image
gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
```

## Using .detectMultiScale() classifier

In [26]:

```
# The next step is to use the face_cascade classifier. I'll let you go explore the
# docs if you
# would like to, but the norm is to use the detectMultiScale() function. This function
# returns
# a list of objects as rectangles. The first parameter is an ndarray of the image.
faces = face_cascade.detectMultiScale(gray)
# And lets just print those faces out to the screen
faces
```

Out[26]:

```
array([[158, 75, 176, 176]], dtype=int32)
```

In [27]:

```
faces.tolist()[0]
```

Out[27]:

```
[158, 75, 176, 176]
```

## Use PILLOW to Draw a Rectangle around the face

The resulting rectangles are in the format of **(x,y,w,h)** where x and y denote the upper left hand point for the image and the width and height represent the bounding box. We know how to handle this in PIL

In [28]:

```
from PIL import Image


# Lets create a PIL image object
pil_img=Image.fromarray(gray,mode="L")

# Now lets bring in our drawing object
from PIL import ImageDraw
# And lets create our drawing context
drawing=ImageDraw.Draw(pil_img)

# Now lets pull the rectangle out of the faces object
rec=faces.tolist()[0]

# Now we just draw a rectangle around the bounds
drawing.rectangle(rec, outline="white")


# And display
display(pil_img)
```



So, not quite what we were looking for. What do you think went wrong? Well, a quick double check of the docs and it is apparent that OpenCV is return the coordinates as  $(x, y, w, h)$  , while PIL.ImageDraw is looking for  $(x1, y1, x2, y2)$  .

In [29]:

```
pil_img=Image.fromarray(gray,mode="L")  
# Setup our drawing context  
drawing=ImageDraw.Draw(pil_img)  
# And draw the new box  
drawing.rectangle((rec[0],rec[1],rec[0]+rec[2],rec[1]+rec[3]), outline="white")  
# And display  
display(pil_img)
```



**Let's try this with a different set of data**

In [30]:

```
img = cv.imread('readonly/msi_recruitment.gif')
# And lets take a look at that image
display(Image.fromarray(img))
```

```
-----
-----
AttributeError                                Traceback (most recent call
1 last)
<ipython-input-30-a6d9e3885cc1> in <module>
      4 img = cv.imread('readonly/msi_recruitment.gif')
      5 # And lets take a look at that image
----> 6 display(Image.fromarray(img))

/opt/conda/lib/python3.7/site-packages/PIL/Image.py in fromarray(obj
j, mode)
    2506     .. versionadded:: 1.1.6
    2507     """
-> 2508     arr = obj.__array_interface__
    2509     shape = arr['shape']
    2510     ndim = len(shape)

AttributeError: 'NoneType' object has no attribute '__array_interfac
e__'
```

Whoa, what's that error about? It looks like there is an error on a line deep within the PIL Image.py file, and it is trying to call an internal **private member** called `_arrayinterface` on the `img` object, but this object is `None`

It turns out that the root of this error is that **OpenCV can't work with Gif images**. This is kind of a pain and unfortunate. But we know how to fix that right? One way is that we could just open this in PIL and then save it as a png, then open that in open cv.

In [31]:

```
# Lets use PIL to open our image
pil_img=Image.open('readonly/msi_recruitment.gif')
# now lets convert it to greyscale for opencv, and get the bytestream
open_cv_version=pil_img.convert("L")
# now lets just write that to a file
open_cv_version.save("msi_recruitment.png")
```

In [32]:

```
# Ok, now that the conversion of format is done, lets try reading this back into
opencv
cv_img=cv.imread('msi_recruitment.png')
# We don't need to color convert this, because we saved it as grayscale
# lets try and detect faces in that image
faces = face_cascade.detectMultiScale(cv_img)

# Now, we still have our PIL color version in a gif
pil_img=Image.open('readonly/msi_recruitment.gif')
# Set our drawing context
drawing=ImageDraw.Draw(pil_img)

# For each item in faces, lets surround it with a red box
for x,y,w,h in faces:
    # That might be new syntax for you! Recall that faces is a list of rectangle
    s in (x,y,w,h)
    # format, that is, a list of lists. Instead of having to do an iteration and
    then manually
    # pull out each item, we can use tuple unpacking to pull out individual item
    s in the sublist
    # directly to variables. A really nice python feature
    #
    # Now we just need to draw our box
    drawing.rectangle((x,y,x+w,y+h), outline="white")
display(pil_img)
```





In [33]:

What happened here!? We see that we have detected faces, and that we have drawn boxes around those faces on the image, but that the colors have gone *all* weird! This, it turns out, has to do *with* color limitations *for* gif images. In short, a gif image has a very limited number of colors. This *is* called a color palette after the palette artists use to mix paints. For gifs the palette can only be 256 colors -- but they can be *\*any\** 56 colors. When a new color *is* introduced, *is* has to take the space of an old color. In this case, PIL adds white to the palette but doesn't know which color to replace and thus messes up the image.

```
# Who knew there was so much to learn about image formats? We can see what mode the image is in with the .mode attribute  
pil_img.mode
```

Out[33]:

'P'

In [34]:


```
# We can see a list of modes in the PILLOW documentation, and they correspond with the color spaces we have been using. For the moment though, lets change back to RGB, which represents color as a three byte tuple instead of in a palette. Lets read in the image  
pil_img=Image.open('readonly/msi_recruitment.gif')  
# Lets convert it to RGB mode  
pil_img = pil_img.convert("RGB")  
# And lets print out the mode  
pil_img.mode
```

Out[34]:

'RGB'

In [35]:

```
# Ok, now lets go back to drawing rectangles. Lets get our drawing object
drawing=ImageDraw.Draw(pil_img)
# And iterate through the faces sequence, tuple unpacking as we go
for x,y,w,h in faces:
    # And remember this is width and height so we have to add those appropriately.
    drawing.rectangle((x,y,x+w,y+h), outline="white")
display(pil_img)
```



In [ ]:

```
# Awesome! We managed to detect a bunch of faces in that image. Looks like we have missed
# four faces. In the machine learning world we would call these false negatives - something
# which the machine thought was not a face (so a negative), but that it was incorrect on.
# Consequently, we would call the actual faces that were detected as true positives -
# something that the machine thought was a face and it was correct on. This leaves us with
# false positives - something the machine thought was a face but it wasn't. We see there are
# two of these in the image, picking up shadow patterns or textures in shirts and matching
# them with the haarcascades. Finally, we have true negatives, or the set of all possible
# rectangles the machine learning classifier could consider where it correctly indicated that
# the result was not a face. In this case there are many many true negatives.
```



In [36]:

```
# There are a few ways we could try and improve this, and really, it requires a
# lot of
# experimentation to find good values for a given image. First, lets create a fu
# nction
# which will plot rectangles for us over the image
def show_rects(faces):
    # Lets read in our gif and convert it
    pil_img=Image.open('readonly/msi_recruitment.gif').convert("RGB")
    # Set our drawing context
    drawing=ImageDraw.Draw(pil_img)
    # And plot all of the rectangles in faces
    for x,y,w,h in faces:
        drawing.rectangle((x,y,x+w,y+h), outline="white")
    # Finally lets display this
    display(pil_img)
```

In [38]:

```
# Ok, first up, we could try and binarize this image. It turns out that opencv h
# as a built in
# binarization function called threshold(). You simply pass in the image, the mi
# dpoint, and
# the maximum value, as well as a flag which indicates whether the threshold sho
# uld be
# binary or something else. Lets try this.
cv_img_bin=cv.threshold(img,120,255,cv.THRESH_BINARY)[1] # returns a list, we wa
# nt the second value
# Now do the actual face detection
faces = face_cascade.detectMultiScale(cv_img_bin)
# Now lets see the results
show_rects(faces)
```



In [ ]:

```
# That's kind of interesting. Not better, but we do see that there is one false  
positive  
# towards the bottom, where the classifier detected the sunglasses as eyes and t  
he dark shadow  
# line below as a mouth.  
#  
# If you're following in the notebook with this video, why don't you pause thing  
s and try a  
# few different parameters for the thresholding value?
```

In [39]:

```
# The detectMultiScale() function from OpenCV also has a couple of parameters. The first of  
# these is the scale factor. The scale factor changes the size of rectangles which are  
# considered against the model, that is, the haarcascades XML file. You can think of it as if  
# it were changing the size of the rectangles which are on the screen.  
#  
# Lets experiment with the scale factor. Usually it's a small value, lets try 1.05  
faces = face_cascade.detectMultiScale(cv_img,1.05)  
# Show those results  
show_rects(faces)  
# Now lets also try 1.15  
faces = face_cascade.detectMultiScale(cv_img,1.15)  
# Show those results  
show_rects(faces)  
# Finally lets also try 1.25  
faces = face_cascade.detectMultiScale(cv_img,1.25)  
# Show those results  
show_rects(faces)
```



## Master of Science in Information



## Master of Science in Information







## Master of Science in Information



In [ ]:

```
# We can see that as we change the scale factor we change the number of true and
# false positives and negatives. With the scale set to 1.05, we have 7 true posi
tives,
# which are correctly identified faces, and 3 false negatives, which are faces w
hich
# are there but not detected, and 3 false positives, where are non-faces which
# opencv thinks are faces. When we change this to 1.15 we lose the false positiv
es but
# also lose one of the true positives, the person to the right wearing a hat. An
d
# when we change this to 1.25 we lost more true positives as well.
#
# This is actually a really interesting phenomena in machine learning and artifi
cial
# intelligence. There is a trade off between not only how accurate a model is, b
ut how
# the inaccuracy actually happens. Which of these three models do you think is b
est?
```

In [40]:

```
# Well, the answer to that question is really, "it depends". It depends why you
# are trying
# to detect faces, and what you are going to do with them. If you think these is
# are interesting, you might want to check out the Applied Data Science with Python
# specialization Michigan offers on Coursera.
#
# Ok, beyond an opportunity to advertise, did you notice anything else that happened when
# we changed the scale factor? It's subtle, but the speed at which the processing ran
# took longer at smaller scale factors. This is because more subimages are being considered
# for these scales. This could also affect which method we might use.
#
# Jupyter has nice support for timing commands. You might have seen this before, a line
# that starts with a percentage sign in jupyter is called a "magic function". This isn't
# normal python - it's actually a shorthand way of writing a function which Jupyter
# has predefined. It looks a lot like the decorators we talked about in a previous
# lecture, but the magic functions were around long before decorators were part of the
# python language. One of the built-in magic functions in jupyter is called timeit, and this
# repeats a piece of python ten times (by default) and tells you the average speed it
# took to complete.
#
# Lets time the speed of detectmultiscale when using a scale of 1.05
%timeit face_cascade.detectMultiScale(cv_img,1.05)
```

1.54 s ± 102 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

In [41]:

```
# Ok, now lets compare that to the speed at scale = 1.15
%timeit face_cascade.detectMultiScale(cv_img,1.15)
```

700 ms ± 49.4 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

In [ ]:

```
# You can see that this is a dramatic difference, roughly two and a half times s  
lower  
# when using the smaller scale!  
#  
# This wraps up our discussion of detecting faces in opencv. You'll see that, li  
ke OCR, this  
# is not a foolproof process. But we can build on the work others have done in m  
achine learning  
# and leverage powerful libraries to bring us closer to building a turn key pyth  
on-based  
# solution. Remember that the detection mechanism isn't specific to faces, tha  
t's just the  
# haarcascades training data we used. On the web you'll be able to find other tr  
aining data  
# to detect other objects, including eyes, animals, and so forth.
```

## More Jupyter Widgets

In [42]:

```
# One of the nice things about using the Jupyter notebook systems is that there  
is a  
# rich set of contributed plugins that seek to extend this system. In this lectu  
re I  
# want to introduce you to one such plugin, call ipy web rtc. Webrtc is a fairly  
new  
# protocol for real time communication on the web. Yup, I'm talking about chatti  
ng.  
# The widget brings this to the Jupyter notebook system. Lets take a look.  
#  
# First, lets import from this library two different classes which we'll use in  
a  
# demo, one for the camera and one for images.  
from ipywebrtc import CameraStream, ImageRecorder  
# Then lets take a look at the camera stream object  
help(CameraStream)
```



Help on class CameraStream in module ipywebrtc.webrtc:

```
class CameraStream(MediaStream)
|   CameraStream(*args, **kwargs)
|
|   Represents a media source by a camera/webcam/microphone using
|   getUserMedia. See
|   https://developer.mozilla.org/en-US/docs/Web/API/MediaDevices/ge
tUserMedia
|   for more detail.
|   The constraints trait can be set to specify constraints for the
camera or
|   microphone, which is described in the documentation of getUserMe
dia, such
|   as in the link above,
|   Two convenience methods are available to easily get access to the
'front'
|   and 'back' camera, when present
|
|   >>> CameraStream.facing_user(audio=False)
|   >>> CameraStream.facing_environment(audio=False)
|
|   Method resolution order:
|       CameraStream
|       MediaStream
|       ipywidgets.widgets.domwidget.DOMWidget
|       ipywidgets.widgets.widget.Widget
|       ipywidgets.widgets.widget.LoggingHasTraits
|       traitlets.traitlets.HasTraits
|       traitlets.traitlets.HasDescriptors
|       builtins.object
|
|   Class methods defined here:
|
|       facing_environment(audio=True, **kwargs) from traitlets.traitlet
s.MetaHasTraits
|       Convenience method to get the camera facing the environment
(often the back)
|
|       Parameters
|       -----
|       audio: bool
|           Capture audio or not
|       kwargs:
|           Extra keyword arguments passed to the `CameraStream`
|
|       facing_user(audio=True, **kwargs) from traitlets.traitlets.MetaH
asTraits
|       Convenience method to get the camera facing the user (often
front)
|
|       Parameters
|       -----
|       audio: bool
|           Capture audio or not
|       kwargs:
|           Extra keyword arguments passed to the `CameraStream`
|
|       -----
|
|   Data descriptors defined here:
```

constraints

Constraints for the camera, see <https://developer.mozilla.org/en-US/docs/Web/API/MediaDevices/getUserMedia> for details.

-----  
Methods inherited from `ipywidgets.widgets.domwidget.DOMWidget`:

`add_class(self, className)`

Adds a class to the top level element of the widget.

Doesn't add the class if it already exists.

`remove_class(self, className)`

Removes a class from the top level element of the widget.

Doesn't remove the class if it doesn't exist.

-----  
Data descriptors inherited from `ipywidgets.widgets.domwidget.DOMWidget`:

`layout`

An instance trait which coerces a dict to an instance.

This lets the instance be specified as a dict, which is used to initialize the instance.

Also, we default to a trivial instance, even if `args` and `kwargs` is not specified.

-----  
Methods inherited from `ipywidgets.widgets.widget.Widget`:

`__del__(self)`

Object disposal

`__init__(self, **kwargs)`

Public constructor

`__repr__(self)`

Return `repr(self)`.

`add_traits(self, **traits)`

Dynamically add trait attributes to the Widget.

`close(self)`

Close method.

Closes the underlying comm.

When the comm is closed, all of the widget views are automatically removed from the front-end.

`get_state(self, key=None, drop_defaults=False)`

Gets the widget state, or a piece of it.

```

Parameters
-----
key : unicode or iterable (optional)
    A single property's name or iterable of property names t
o get.

Returns
-----
state : dict of states
metadata : dict
    metadata for each field: {key: metadata}

get_view_spec(self)

hold_sync(self)
    Hold syncing any state until the outermost context manager e
xits

notify_change(self, change)
    Called when a property has changed.

on_displayed(self, callback, remove=False)
    (Un)Register a widget displayed callback.

Parameters
-----
callback: method handler
    Must have a signature of::

        callback(widget, **kwargs)

    kwargs from display are passed through without modificat
ion.

remove: bool
    True if the callback should be unregistered.

on_msg(self, callback, remove=False)
    (Un)Register a custom msg receive callback.

Parameters
-----
callback: callable
    callback will be passed three arguments when a message a
rrives::

        callback(widget, content, buffers)

remove: bool
    True if the callback should be unregistered.

open(self)
    Open a comm to the frontend if one isn't already open.

send(self, content, buffers=None)
    Sends a custom msg to the widget model in the front-end.

Parameters
-----
content : dict
    Content of the message to send.
buffers : list of binary buffers

```

Binary buffers to send with message

`send_state(self, key=None)`

Sends the widget state, or a piece of it, to the front-end, if it exists.

Parameters

-----

`key` : unicode, or iterable (optional)

A single property's name or iterable of property names to sync with the front-end.

`set_state(self, sync_data)`

Called when a state is received from the front-end.

-----

Class methods inherited from `ipywidgets.widgets.widget.Widget`:

`close_all()` from `traitlets.traitlets.MetaHasTraits`

-----

Static methods inherited from `ipywidgets.widgets.widget.Widget`:

`get_manager_state(drop_defaults=False, widgets=None)`

Returns the full state for a widget manager for embedding

:param drop\_defaults: when True, it will not include default value

:param widgets: list with widgets to include in the state (or all widgets when None)

:return:

`handle_comm_opened(comm, msg)`

Static method, called when a widget is constructed.

`on_widget_constructed(callback)`

Registers a callback to be called when a widget is constructed.

The callback must have the following signature:

`callback(widget)`

-----

Data descriptors inherited from `ipywidgets.widgets.widget.Widget`:

`comm`

A trait whose value must be an instance of a specified class.

The value can also be an instance of a subclass of the specified class.

Subclasses can declare default classes by overriding the `klass` attribute

`keys`

The traits which are synced.

model\_id  
Gets the model id of this widget.  
  
If a Comm doesn't exist yet, a Comm will be created automagically.

-----  
Data and other attributes inherited from ipywidgets.widget.Widget:  
  
widget\_types = <ipywidgets.widget.WidgetRegistry object>  
  
widgets = {}  
  
-----

-----  
Data descriptors inherited from ipywidgets.widget.LoggingHasTraits:  
  
log  
A trait whose value must be an instance of a specified class.  
  
The value can also be an instance of a subclass of the specified class.  
  
Subclasses can declare default classes by overriding the klass attribute

-----  
Methods inherited from traitlets.traitlets.HasTraits:  
  
\_\_getstate\_\_(self)  
  
\_\_setstate\_\_(self, state)  
  
has\_trait(self, name)  
Returns True if the object has a trait with the specified name.  
  
hold\_trait\_notifications(self)  
Context manager for bundling trait change notifications and cross validation.  
  
Use this when doing multiple trait assignments (init, config), to avoid race conditions in trait notifiers requesting other trait values.  
  
All trait notifications will fire after all values have been assigned.  
  
observe(self, handler, names=traitlets.All, type='change')  
Setup a handler to be called when a trait changes.  
  
This is used to setup dynamic notifications of trait changes.

```

Parameters
-----
handler : callable
    A callable that is called when a trait changes. Its
    signature should be ``handler(change)``, where ``change``
is a dictionary. The change dictionary at least holds a 'type'
key.
    * ``type``: the type of notification.
    Other keys may be passed depending on the value of 'type'. In the
case where type is 'change', we also have the following
keys:
    * ``owner`` : the HasTraits instance
    * ``old`` : the old value of the modified trait attribute
    * ``new`` : the new value of the modified trait attribute
    * ``name`` : the name of the modified trait attribute.
names : list, str, All
    If names is All, the handler will apply to all traits.
If a list
    of str, handler will apply to all names in the list. If
a
    str, the handler will apply just to that name.
type : str, All (default: 'change')
    The type of notification to filter by. If equal to All,
then all
    notifications are passed to the observe handler.

on_trait_change(self, handler=None, name=None, remove=False)
    DEPRECATED: Setup a handler to be called when a trait change
s.
    This is used to setup dynamic notifications of trait change
s.
    Static handlers can be created by creating methods on a HasT
raits
    subclass with the naming convention '_[traitname]_changed'.
Thus,
    to create static handler for the trait 'a', create the metho
d
    _a_changed(self, name, old, new) (fewer arguments can be use
d, see
    below).
    If `remove` is True and `handler` is not specified, all chan
ge
    handlers for the specified name are uninstalled.

Parameters
-----
handler : callable, None
    A callable that is called when a trait changes. Its
    signature can be handler(), handler(name), handler(name,
new),
    handler(name, old, new), or handler(name, old, new, sel
f).
    name : list, str, None
    If None, the handler will apply to all traits. If a lis

```

```

t
|         of str, handler will apply to all names in the list.  If
a
|         str, the handler will apply just to that name.
|         remove : bool
|         If False (the default), then install the handler.  If Tr
ue
|         then unintall it.
|
|         set_trait(self, name, value)
|         Forcibly sets trait attribute, including read-only attribute
s.
|
|         setup_instance(self, *args, **kwargs)
|         This is called **before** self.__init__ is called.
|
|         trait_metadata(self, traitname, key, default=None)
|         Get metadata values for trait by key.
|
|         trait_names(self, **metadata)
|         Get a list of all the names of this class' traits.
|
|         traits(self, **metadata)
|         Get a ``dict`` of all the traits of this class.  The diction
ary
|         is keyed on the name and the values are the TraitType object
s.
|
|         The TraitTypes returned don't know anything about the values
|         that the various HasTrait's instances are holding.
|
|         The metadata kwargs allow functions to be passed in which
|         filter traits based on metadata values.  The functions shoul
d
|         take a single value as an argument and return a boolean.  If
|         any function returns False, then the trait is not included i
n
|         the output.  If a metadata key doesn't exist, None will be p
assed
|         to the function.
|
|         unobserve(self, handler, names=traitlets.All, type='change')
|         Remove a trait change handler.
|
|         This is used to unregister handlers to trait change notifica
tions.
|
|         Parameters
|         -----
|         handler : callable
|             The callable called when a trait attribute changes.
|         names : list, str, All (default: All)
|             The names of the traits for which the specified handler
should be
|             uninstalled.  If names is All, the specified handler is u
ninstalled
|             from the list of notifiers corresponding to all changes.
|         type : str or All (default: 'change')
|             The type of notification to filter by.  If All, the speci
fied handler
|             is uninstalled from the list of notifiers corresponding

```

to all types.

```
unobserve_all(self, name=traitlets.All)
```

Remove trait change handlers of any type for the specified name.

If name is not specified, removes all trait notifiers.

-----

Class methods inherited from traitlets.traitlets.HasTraits:

```
class_own_trait_events(name) from traitlets.traitlets.MetaHasTra
```

its

Get a dict of all event handlers defined on this class, not a parent.

Works like ``event\_handlers``, except for excluding traits from parents.

```
class_own_traits(**metadata) from traitlets.traitlets.MetaHasTra
```

its

Get a dict of all the traitlets defined on this class, not a parent.

Works like `class\_traits`, except for excluding traits from parents.

```
class_trait_names(**metadata) from traitlets.traitlets.MetaHasTra
```

aits

Get a list of all the names of this class' traits.

This method is just like the :meth:`trait\_names` method, but is unbound.

```
class_traits(**metadata) from traitlets.traitlets.MetaHasTraits
```

ary

Get a ``dict`` of all the traits of this class. The dictionary

s.

is keyed on the name and the values are the TraitType objects.

s.

s.

s.

nbound.

The TraitTypes returned don't know anything about the values that the various HasTrait's instances are holding.

The metadata kwargs allow functions to be passed in which filter traits based on metadata values. The functions should

d

take a single value as an argument and return a boolean. If any function returns False, then the trait is not included in

n

the output. If a metadata key doesn't exist, None will be passed

assed

to the function.

```
trait_events(name=None) from traitlets.traitlets.MetaHasTraits
```

Get a ``dict`` of all the event handlers of this class.

Parameters

-----



```

    name: str (default: None)
        The name of a trait of this class. If name is ``None`` t
hen all
        the event handlers of this class will be returned instea
d.

    Returns
    -----
    The event handlers associated with a trait name, or all even
t handlers.

-----

Data descriptors inherited from traitlets.traitlets.HasTraits:

    cross_validation_lock
        A contextmanager for running a block with our cross validati
on lock set
        to True.

        At the end of the block, the lock's value is restored to its
value
        prior to entering the block.

-----

Static methods inherited from traitlets.traitlets.HasDescriptor
s:

    __new__(cls, *args, **kwargs)
        Create and return a new object. See help(type) for accurate
signature.

-----

Data descriptors inherited from traitlets.traitlets.HasDescripto
rs:

    __dict__
        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)

```

In [43]:

```
# We see from the docs that it's east to get a camera facing the user, and we can have  
# the audio on or off. We don't need audio for this demo, so lets create a new camera  
# instance  
camera = CameraStream.facing_user(audio=False)  
# The next object we want to look at is the ImageRecorder  
help(ImageRecorder)
```

Help on class ImageRecorder in module ipywebrtc.webrtc:

```
class ImageRecorder(Recorder)
|   ImageRecorder(*args, **kwargs)
|
|   Creates a recorder which allows to grab an Image from a MediaStr
eam widget.
|
|   Method resolution order:
|       ImageRecorder
|       Recorder
|       ipywidgets.widgets.domwidget.DOMWidget
|       ipywidgets.widgets.widget.Widget
|       ipywidgets.widgets.widget.LoggingHasTraits
|       traitlets.traitlets.HasTraits
|       traitlets.traitlets.HasDescriptors
|       builtins.object
|
|   Methods defined here:
|
|   __init__(self, format='png', filename='record', recording=False,
autosave=False, **kwargs)
|       Public constructor
|
|   save(self, filename=None)
|       Save the image to a file, if no filename is given it is base
d on the filename trait and the format.
|
|       >>> recorder = ImageRecorder(filename='test', format='png')
|       >>> ...
|       >>> recorder.save() # will save to test.png
|       >>> recorder.save('foo') # will save to foo.png
|       >>> recorder.save('foo.dat') # will save to foo.dat
|
| -----
|
| Data descriptors defined here:
|
| format
|     The format of the image.
|
| image
|     A trait whose value must be an instance of a specified clas
s.
|
|     The value can also be an instance of a subclass of the speci
fied class.
|
|     Subclasses can declare default classes by overriding the kla
ss attribute
|
| -----
|
| Methods inherited from Recorder:
|
| download(self)
|     Download the recording (usually a popup appears in the brows
er)
|
| -----
|
| -----
```

Data descriptors inherited from Recorder:

autosave

If true, will save the data to a file once the recording is finished (based on filename and format)

filename

The filename used for downloading or auto saving.

recording

(boolean) Indicator and controller of the recorder state, i.e. putting the value to True will start recording.

stream

An instance of :class:`MediaStream` that is the source for recording.

-----  
Methods inherited from ipywidgets.widgets.domwidget.DOMWidget:

add\_class(self, className)

Adds a class to the top level element of the widget.

Doesn't add the class if it already exists.

remove\_class(self, className)

Removes a class from the top level element of the widget.

Doesn't remove the class if it doesn't exist.

-----  
Data descriptors inherited from ipywidgets.widgets.domwidget.DOMWidget:

layout

An instance trait which coerces a dict to an instance.

This lets the instance be specified as a dict, which is used to initialize the instance.

Also, we default to a trivial instance, even if args and kwargs is not specified.

-----  
Methods inherited from ipywidgets.widget.Widget:

\_\_del\_\_(self)

Object disposal

\_\_repr\_\_(self)

Return repr(self).

add\_traits(self, \*\*traits)

Dynamically add trait attributes to the Widget.

close(self)

Close method.

```

        Closes the underlying comm.
        When the comm is closed, all of the widget views are automat
ically
        removed from the front-end.

    get_state(self, key=None, drop_defaults=False)
        Gets the widget state, or a piece of it.

        Parameters
        -----
        key : unicode or iterable (optional)
            A single property's name or iterable of property names t
o get.

        Returns
        -----
        state : dict of states
        metadata : dict
            metadata for each field: {key: metadata}

    get_view_spec(self)

    hold_sync(self)
        Hold syncing any state until the outermost context manager e
xits

    notify_change(self, change)
        Called when a property has changed.

    on_displayed(self, callback, remove=False)
        (Un)Register a widget displayed callback.

        Parameters
        -----
        callback: method handler
            Must have a signature of::

                callback(widget, **kwargs)

            kwargs from display are passed through without modificat
ion.

        remove: bool
            True if the callback should be unregistered.

    on_msg(self, callback, remove=False)
        (Un)Register a custom msg receive callback.

        Parameters
        -----
        callback: callable
            callback will be passed three arguments when a message a
rrives::

                callback(widget, content, buffers)

        remove: bool
            True if the callback should be unregistered.

    open(self)
        Open a comm to the frontend if one isn't already open.

```

```

send(self, content, buffers=None)
    Sends a custom msg to the widget model in the front-end.

    Parameters
    -----
    content : dict
        Content of the message to send.
    buffers : list of binary buffers
        Binary buffers to send with message

send_state(self, key=None)
    Sends the widget state, or a piece of it, to the front-end,
if it exists.

    Parameters
    -----
    key : unicode, or iterable (optional)
        A single property's name or iterable of property names t
o sync with the front-end.

set_state(self, sync_data)
    Called when a state is received from the front-end.

-----

Class methods inherited from ipywidgets.widgets.widget.Widget:

close_all() from traitlets.traitlets.MetaHasTraits

-----

Static methods inherited from ipywidgets.widgets.widget.Widget:

get_manager_state(drop_defaults=False, widgets=None)
    Returns the full state for a widget manager for embedding

    :param drop_defaults: when True, it will not include default
value
    :param widgets: list with widgets to include in the state (o
r all widgets when None)
    :return:

handle_comm_opened(comm, msg)
    Static method, called when a widget is constructed.

on_widget_constructed(callback)
    Registers a callback to be called when a widget is construct
ed.

    The callback must have the following signature:
    callback(widget)

-----

Data descriptors inherited from ipywidgets.widgets.widget.Widge
t:

comm
    A trait whose value must be an instance of a specified clas
s.

```

The value can also be an instance of a subclass of the specified class.

Subclasses can declare default classes by overriding the `klass` attribute

`keys`

The traits which are synced.

`model_id`

Gets the model id of this widget.

If a `Comm` doesn't exist yet, a `Comm` will be created automatically.

Data and other attributes inherited from `ipywidgets.widget.Widget`:

`widget_types = <ipywidgets.widget.WidgetRegistry object>`

`widgets = {'cc401cd050ea4415bc5e87f7e3a0e0dd': CameraStream(Constraint...`

Data descriptors inherited from `ipywidgets.widget.LoggingHasTraits`:

`log`

A trait whose value must be an instance of a specified class.

The value can also be an instance of a subclass of the specified class.

Subclasses can declare default classes by overriding the `klass` attribute

Methods inherited from `traitlets.traitlets.HasTraits`:

`__getstate__(self)`

`__setstate__(self, state)`

`has_trait(self, name)`

Returns True if the object has a trait with the specified name.

`hold_trait_notifications(self)`

Context manager for bundling trait change notifications and validation.

Use this when doing multiple trait assignments (`init`, `config`), to avoid race conditions in trait notifiers requesting other trait va

lues.

All trait notifications will fire after all values have been assigned.

```
observe(self, handler, names=traitlets.All, type='change')
```

Setup a handler to be called when a trait changes.

This is used to setup dynamic notifications of trait change

s.

Parameters

-----

handler : callable

A callable that is called when a trait changes. Its signature should be ``handler(change)``, where ``change``

is a

dictionary. The change dictionary at least holds a 'type

e' key.

\* ``type``: the type of notification.

Other keys may be passed depending on the value of 'type

e'. In the

case where type is 'change', we also have the following

keys:

\* ``owner`` : the HasTraits instance

\* ``old`` : the old value of the modified trait attribute

e

\* ``new`` : the new value of the modified trait attribute

e

\* ``name`` : the name of the modified trait attribute.

names : list, str, All

If names is All, the handler will apply to all traits.

If a list

of str, handler will apply to all names in the list. If

a

str, the handler will apply just to that name.

type : str, All (default: 'change')

The type of notification to filter by. If equal to All,

then all

notifications are passed to the observe handler.

```
on_trait_change(self, handler=None, name=None, remove=False)
```

DEPRECATED: Setup a handler to be called when a trait change

s.

This is used to setup dynamic notifications of trait change

s.

Static handlers can be created by creating methods on a HasT

raits

subclass with the naming convention '\_[traitname]\_changed'.

Thus,

to create static handler for the trait 'a', create the metho

d

\_a\_changed(self, name, old, new) (fewer arguments can be use

d, see

below).

If ``remove`` is True and ``handler`` is not specified, all chan

ge

handlers for the specified name are uninstalled.



```

Parameters
-----
handler : callable, None
    A callable that is called when a trait changes. Its
    signature can be handler(), handler(name), handler(name,
new),
    handler(name, old, new), or handler(name, old, new, sel
f).
name : list, str, None
    If None, the handler will apply to all traits. If a lis
t
    of str, handler will apply to all names in the list. If
a
    str, the handler will apply just to that name.
remove : bool
    If False (the default), then install the handler. If Tr
ue
    then uninstall it.

set_trait(self, name, value)
    Forcibly sets trait attribute, including read-only attribute
s.

setup_instance(self, *args, **kwargs)
    This is called before self.__init__ is called.

trait_metadata(self, traitname, key, default=None)
    Get metadata values for trait by key.

trait_names(self, **metadata)
    Get a list of all the names of this class' traits.

traits(self, **metadata)
    Get a ``dict`` of all the traits of this class. The diction
ary
    is keyed on the name and the values are the TraitType object
s.

The TraitTypes returned don't know anything about the values
that the various HasTrait's instances are holding.

The metadata kwargs allow functions to be passed in which
filter traits based on metadata values. The functions shoul
d
    take a single value as an argument and return a boolean. If
    any function returns False, then the trait is not included i
n
    the output. If a metadata key doesn't exist, None will be p
assed
    to the function.

unobserve(self, handler, names=traitlets.All, type='change')
    Remove a trait change handler.

This is used to unregister handlers to trait change notifica
tions.

Parameters
-----
handler : callable
    The callable called when a trait attribute changes.

```

```

|         names : list, str, All (default: All)
|             The names of the traits for which the specified handler
should be
|             uninstalled. If names is All, the specified handler is u
ninstalled
|             from the list of notifiers corresponding to all changes.
|         type : str or All (default: 'change')
|             The type of notification to filter by. If All, the speci
fied handler
|             is uninstalled from the list of notifiers corresponding
to all types.
|
|         unobserve_all(self, name=traitlets.All)
|             Remove trait change handlers of any type for the specified n
ame.
|             If name is not specified, removes all trait notifiers.
|
|         -----

```

```

|-----
| Class methods inherited from traitlets.traitlets.HasTraits:
|
|         class_own_trait_events(name) from traitlets.traitlets.MetaHasTra
its
|             Get a dict of all event handlers defined on this class, not
a parent.
|
|             Works like ``event_handlers``, except for excluding traits f
rom parents.
|
|         class_own_traits(**metadata) from traitlets.traitlets.MetaHasTra
its
|             Get a dict of all the traitlets defined on this class, not a
parent.
|
|             Works like `class_traits`, except for excluding traits from
parents.
|
|         class_trait_names(**metadata) from traitlets.traitlets.MetaHasTr
aits
|             Get a list of all the names of this class' traits.
|
|             This method is just like the :meth:`trait_names` method,
but is unbound.
|
|         class_traits(**metadata) from traitlets.traitlets.MetaHasTraits
|             Get a ``dict`` of all the traits of this class. The diction
ary
|             is keyed on the name and the values are the TraitType object
s.
|
|             This method is just like the :meth:`traits` method, but is u
nbound.
|
|             The TraitTypes returned don't know anything about the values
that the various HasTrait's instances are holding.
|
|             The metadata kwargs allow functions to be passed in which
filter traits based on metadata values. The functions shoul
d
|             take a single value as an argument and return a boolean. If
any function returns False, then the trait is not included i

```

```

n
|         the output.  If a metadata key doesn't exist, None will be p
assed
|         to the function.
|
|         trait_events(name=None) from traitlets.traitlets.MetaHasTraits
|         Get a ``dict`` of all the event handlers of this class.
|
|         Parameters
|         -----
|         name: str (default: None)
|             The name of a trait of this class. If name is ``None`` t
hen all
|             the event handlers of this class will be returned instea
d.
|
|         Returns
|         -----
|         The event handlers associated with a trait name, or all even
t handlers.
|
|         -----
|
|-----
|         Data descriptors inherited from traitlets.traitlets.HasTraits:
|
|         cross_validation_lock
|             A contextmanager for running a block with our cross validati
on lock set
|             to True.
|
|             At the end of the block, the lock's value is restored to its
value
|             prior to entering the block.
|
|         -----
|
|-----
|         Static methods inherited from traitlets.traitlets.HasDescriptor
s:
|
|         __new__(cls, *args, **kwargs)
|             Create and return a new object.  See help(type) for accurate
signature.
|
|         -----
|
|-----
|         Data descriptors inherited from traitlets.traitlets.HasDescripto
rs:
|
|         __dict__
|             dictionary for instance variables (if defined)
|
|         __weakref__
|             list of weak references to the object (if defined)

```

In [46]:

```
# The image recorder lets us actually grab images from the camera stream. There
# are features
# for downloading and using the image as well. We see that the default format is
# a png file.
# Lets hook up the ImageRecorder to our stream
image_recorder = ImageRecorder(stream=camera)
# Now, the docs are a little unclear how to use this within Jupyter, but if we c
# all the
# download() function it will actually store the results of the camera which is
# hooked up
# in image_recorder.image. Lets try it out
# First, lets tell the recorder to start capturing data
image_recorder.recording=True
# Now lets download the image
image_recorder.download()
# Then lets inspect the type of the image
type(image_recorder.image)
```

Out[46]:

```
ipywidgets.widgets.widget_media.Image
```

In [47]:

```
# Ok, the object that it stores is an ipywidgets.widgets.widget_media.Image. How
do we do
# something useful with this? Well, an inspection of the object shows that there
is a handy
# value field which actually holds the bytes behind the image. And we know how t
o display
# those.
# Lets import PIL Image
import PIL.Image
# And lets import io
import io
# And now lets create a PIL image from the bytes
img = PIL.Image.open(io.BytesIO(image_recorder.image.value))
# And render it to the screen
display(img)
```

```
-----
-----
OSError                                Traceback (most recent call
last)
<ipython-input-47-bal7cbdcfa82> in <module>
      8 import io
      9 # And now lets create a PIL image from the bytes
--> 10 img = PIL.Image.open(io.BytesIO(image_recorder.image.value))
     11 # And render it to the screen
     12 display(img)

/opt/conda/lib/python3.7/site-packages/PIL/Image.py in open(fp, mod
e)
    2685         warnings.warn(message)
    2686         raise IOError("cannot identify image file %r"
-> 2687                     % (filename if filename else fp))
    2688
    2689 #

OSError: cannot identify image file <_io.BytesIO object at 0x7f132b8
1daf0>
```

In [ ]:

```
# Great, you see a picture! Hopefully you are following along in one of the note
books
# and have been able to try this out for yourself!
#
# What can you do with this? This is a great way to get started with a bit of co
mputer vision.
# You already know how to identify a face in the webcam picture, or try and capt
ure text
# from within the picture. With OpenCV there are any number of other things you
can do, simply
# with a webcam, the Jupyter notebooks, and python!
```